



Forge4Flame (F4F): User Guide

Daniele Baccega, Irene Terrone, Peter Heywood, Robert Chisholm, Paul Richmond, Sandro Gepiro Contaldo, Lorenzo Bosio, Simone Pernice, Marco Beccuti

Contents

1	Introduction	2
2	How to download	2
2.1	Repository	2
2.2	F4F	2
2.2.1	Dependencies	3
2.3	FLAME GPU 2	3
2.4	Dockers	3
2.5	Slurm	4
3	How to run	4
3.1	F4F	4
3.1.1	Without Docker	4
3.1.2	With Docker	4
3.2	FLAME GPU 2	4
3.2.1	Without Docker	4
3.2.2	With Docker	5
3.3	Docker Compose	6
3.4	Slurm	7
3.5	How to reproduce the results	8
4	Forge4Flame application	8
4.1	Homepage	8
4.2	Canvas	9
4.2.1	Floor Plan Import as Background Image Support	11
4.3	Rooms	12
4.4	Agents	14
4.5	Resources	18
4.6	Infection	19
4.7	What-If	21
4.8	Configuration	27
4.9	Run	28
4.10	Settings	30

4.11 Post Processing	30
4.11.1 Uploading simulation	31
4.11.2 2D Simulation Visualisation	32
4.11.3 Contact Matrix	35
5 School	36
6 Link between F4F and FLAME GPU 2	43

1 Introduction

This document serves as a user guide for *Forge4Flame (F4F)*. Specifically, Section 2 provides various instructions for downloading the repository, *F4F*, FLAME GPU 2, the Docker images, and Slurm, while Section 3 explains how to run them in different ways. Section 4 introduces the main features of the application, using an example as a guide. Section 5 provides additional details about the generic middle school environment discussed in the main document. Finally, Section 6 briefly explains the integration between *F4F* and FLAME GPU 2.

2 How to download

This section provides instructions for downloading the repository, the *F4F* R package, FLAME GPU 2, and the Docker images and configuring the HPC environment for Slurm.

2.1 Repository

Cloning the repository is optional, especially for users intending to use the *F4F* package, dockers, or Slurm—see the following sections for more details. The GitHub repository is available at <https://github.com/qBioTurin/FORGE4FLAME>. To clone it, use the following Bash command:

```
git clone --recurse-submodules \
    https://github.com/qBioTurin/FORGE4FLAME.git
```

The main directory, FORGE4FLAME, contains the R Shiny application, while the FORGE4FLAME /inst/FLAMEGPU-FORGE4FLAME directory stores the FLAME GPU 2 template used by *F4F*.

2.2 F4F

To install the *F4F* R package it is possible to use the R package *devtools* [28] or *remotes* [11], through these R commands:

```
if(!requireNamespace("devtools", quietly = TRUE))
    install.packages("devtools")

library(devtools)
install_github("https://github.com/qBioTurin/FORGE4FLAME")
```

or:

```
if(!requireNamespace("remotes", quietly = TRUE))
    install.packages("remotes")

library(remotes)
```

```
install_github("https://github.com/qBioTurin/FORGE4FLAME")
```

2.2.1 Dependencies

F4F is implemented in R Shiny [8]. It requires several CRAN packages to operate: *shinydashboard* [7], *shinyjs* [4], *jsonlite* [16], *dplyr* [29], *shinythemes* [9], *colourpicker* [2], *glue* [13], *readr* [30], *zip* [10], *sortable* [24], *shinyalert* [3], *shinybusy* [15], *shinyBS* [5], *stringr* [27], *ggplot2* [25], *tidyR* [26], *DT* [31], *shiny* [6], *shinyWidgets* [21], *shinyFiles* [19], and *htmltools* [22]. These packages are automatically installed.

Furthermore, it also requires the EBImage [18] and the EBImageExtra [17] packages. The user must install these packages manually using the following R commands:

```
if(!requireNamespace("BiocManager", quietly = TRUE))
  install.packages("BiocManager")

BiocManager::install("EBImage")

if(!requireNamespace("remotes", quietly = TRUE))
  install.packages("remotes")

library(remotes)
install_github("ornelles/EBImageExtra")
```

2.3 FLAME GPU 2

To install FLAME GPU 2 dependencies, refer to the official documentation here.

2.4 Dockers

Users can download the Docker images for both *F4F* and FLAME GPU 2 to avoid any potential dependency-related issues. In this context, the user must have Docker installed on their computer. For more information, refer to this document: <https://docs.docker.com/engine/installation/>. Additionally, the user needs to ensure they have the necessary permissions to run Docker without using sudo. To create the Docker group and add the user on a Unix system, follow these steps:

- Create the docker group:

```
$ sudo groupadd docker
```

- Add user to the docker group:

```
$ sudo usermod -aG docker $USER
```

- Log out and log back in so that group membership is re-evaluated.

Additionally, the user must have the NVIDIA driver (installation instructions can be found here) and the NVIDIA container toolkit (installation instructions can be found here) installed. To download the Docker images, run the following Bash commands:

```
docker pull qbioturin/forge4flame
docker pull qbioturin/flamegpu2
```

An alternative method for downloading Docker images is via *F4F*, provided that the user has installed the R package (refer to Section 2 for installation instructions):

```
library(FORGE4FLAME)
FORGE4FLAME::downloadContainers()
```

2.5 Slurm

To run FLAME GPU 2 simulations on an HPC system, the user must install Slurm on it (more information here). Generally, on HPC systems, Docker cannot be used due to privacy concerns. Therefore, the user must install all necessary FLAME GPU 2 dependencies on the user’s system before running simulations (see Section 2.3 for more details). However, on HPC4AI [1] (more information here), FLAME GPU 2 can also be executed using Docker, thanks to a tool that addresses privacy concerns.

3 How to run

This section describes how to run *F4F* and FLAME GPU 2, both with and without Docker, using the three approaches outlined in the main document (see Figure 2).

3.1 F4F

3.1.1 Without Docker

To run *F4F* without using Docker, run the following R commands after installing the package (see Section 2.2):

```
library(FORGE4FLAME)
FORGE4FLAME::FORGE4FLAME.run()
```

3.1.2 With Docker

To run it using Docker, execute the following Bash command (if running on a server, ensure that port 3838 is exposed and accessible via `http://<server-hostname>:3838`; if running locally, access at `http://localhost:3838`):

```
docker run -p 3838:3838 qbioturin/forge4flame
```

In this case, there is no need to download the *F4F* repository or install the package. In addition, it is not possible to run a FLAME GPU 2 simulation using the *Run* button in the **Run** page—see Section 4.9 for more details.

3.2 FLAME GPU 2

3.2.1 Without Docker

To run FLAME GPU 2 without using Docker, run the following Bash command (inside the `FORGE4FLAME/inst/FLAMEGPU-FORGE4FLAME` directory, after cloning the *F4F* repository):

```

# Executing a single run without using the FLAME GPU 2
# 3D visualisation:

./abm.sh -expdir NameOfTheModel

#####
# Executing a single run using the FLAME GPU 2
# 3D visualisation:

./abm.sh -expdir NameOfTheModel \
-v ON

#####
# Executing n runs without using the FLAME GPU 2
# 3D visualisation:

./abm_ensemble.sh -expdir NameOfTheModel

#####
# visualise the helper:

./abm.sh -h
./abm_ensemble.sh -h

```

In particular, `NameOfTheModel` must be the name of the desired model and must correspond to a directory within `FORGE4FLAME/inst/FLAMEGPU-FORGE4FLAME/resources/f4f` that contains both a JSON file and an RDs file. If the user utilises the *Link the model to FLAME GPU 2* button (see Section 4.10 for more details), the directory with the provided name will be automatically saved in the correct location. Otherwise, if the user downloads the model using the *Save the model* button, they must manually copy the unzipped directory into `FORGE4FLAME/inst/FLAMEGPU-FORGE4FLAME/resources/f4f`. Results will be saved in `FORGE4FLAME/inst/FLAMEGPU-FORGE4FLAME/results/NameOfTheModel`.

3.2.2 With Docker

To run FLAME GPU 2 using Docker, run the following Bash commands:

```

# Executing a single run without using the FLAME GPU 2
# 3D visualisation:

docker run --user $UID:$UID \
--rm \
--gpus all \
--runtime nvidia \
-v /Absolute/Path/To/The/Directory/With/The/Model/
NameOfTheModel:/home/docker/flamegpu2/FLAMEGPU-
FORGE4FLAME/resources/f4f/NameOfTheModel \

```

```

-v $(pwd):/home/docker/flamegpu2/FLAMEGPU-FORGE4FLAME/
    flamegpu2_results \
qbioturin/flamegpu2 /usr/bin/bash -c "/home/docker/
    flamegpu2/FLAMEGPU-FORGE4FLAME/abm.sh -expdir
    NameOfTheModel"

#####
# Executing n runs without using the FLAME GPU 2
# 3D visualisation:

docker run --user $UID:$UID \
    --rm \
    --gpus all \
    --runtime nvidia \
    -v /Absolute/Path/To/The/Directory/With/The/Model/
        NameOfTheModel:/home/docker/flamegpu2/FLAMEGPU-
        FORGE4FLAME/resources/f4f/NameOfTheModel \
    -v $(pwd):/home/docker/flamegpu2/FLAMEGPU-FORGE4FLAME/
        flamegpu2_results \
qbioturin/flamegpu2 /usr/bin/bash -c "/home/docker/
    flamegpu2/FLAMEGPU-FORGE4FLAME/abm_ensemble.sh -expdir
    NameOfTheModel"

```

In this case, there is no need to download the *F4F* repository or install the package. */Absolute/Path/To/The/Directory/With/The/Model/NameOfTheModel* represents the **absolute** path to the local directory, named *NameOfTheModel*, which contains the model to run (the JSON and the RDs files). These files will be saved in a directory named *NameOfTheModel* inside the Docker (in *FLAMEGPU-FORGE4FLAME/resources/f4f*). Results will be saved in a directory named *results/NameOfTheModel* within the current directory. The user must replace the directory name containing the model with *NameOfTheModel*.

3.3 Docker Compose

To use Docker Compose, the user must download the YAML file here. To start both *F4F* and FLAME GPU 2 containers, navigate to the directory containing the YAML file and run the following Bash command (if running on a server, ensure that port 3839 is exposed and accessible via <http://<server-hostname>:3839>; if running locally, access to <http://localhost:3839>):

```
docker compose up -d --build
```

To run a FLAME GPU 2 simulation using Docker Compose, the user must use the **Run** page of *F4F*—see Section 3 for more details. Results will be saved in a directory named *results/NameOfTheModel* within the current directory, where *NameOfTheModel* is the name selected by the user when clicking on the *Run* button in the **Run** page—see Section 4.9 for more details.

To stop the containers, run the following Bash command:

```
docker compose down
```

3.4 Slurm

To execute FLAME GPU 2 simulations on HPC systems using Slurm, assuming the necessary drivers are installed, the repository is cloned, and one or more reservations have been made on the HPC system, a Bash script like the one below can be used (this script, named `run_on_slurm.sh`, is included in the repository):

```
#!/bin/bash

# Define arrays
PARTITIONS=("MachineName1-booked", "MachineName2-booked",
    ... , "MachineNameN-booked")
RESERVATIONS=("Run1", "Run2", ... , "RunN")
GRES=("GPUName1", "GPUName2", ... , "GPUNameN")
MODEL_NAMES=("NameOfTheModel1", "NameOfTheModel2", ... ,
    NameOfTheModelN")

ARRAY_LENGTH=${#PARTITIONS[@]}

for ((i=0; i<ARRAY_LENGTH; i++)); do
    JOB_SCRIPT="job_script$i.sh"

    cat <<EOF > "$JOB_SCRIPT"
#!/bin/bash
#SBATCH --partition=${PARTITIONS[$i]}
#SBATCH --reservation=${RESERVATIONS[$i]}
#SBATCH -N 1
#SBATCH --gres=${GRES[$i]}

./abm_ensemble.sh -expdir ${MODEL_NAMES[$i]}
EOF

    # Submit the job and check if it was successful
    if sbatch "$JOB_SCRIPT"; then
        echo "Submitted: $JOB_SCRIPT"
    else
        echo "Failed to submit: $JOB_SCRIPT" >&2
    fi
done
```

In this script, the user can define a set of partitions, reservations, GPUs, and model names by specifying the four vectors:

- `PARTITIONS` specifies the names of the machines reserved.
- `RESERVATIONS` defines the reservation names chosen by the user at book time.
- `GRES` determines the name of the GPUs.
- `MODEL_NAMES` specifies the names of the models to run, which must correspond to the folder names within `FLAMEGPU-FORGE4FLAME/resources/f4f`.

Specifically, we run the experiments presented in the main document using Slurm on HPC4AI [1] (more information [here](#)).

3.5 How to reproduce the results

To reproduce the results presented in the main document—comparing FLAME GPU 2 and NetLogo using the ITP and the alert scenarios—run the following Bash commands (note that this process may take time, especially for NetLogo):

```
git clone -b School --recurse-submodules \
https://github.com/qBioTurin/FLAMEGPU-FORGE4FLAME.git

cd FLAMEGPU-FORGE4FLAME

# Reproduce results on comparison
# between FLAME GPU 2 and NetLogo:
./reproduce_comparison.sh

# Reproduce results on alert
# scenarios:
./reproduce_alert.sh
```

4 Forge4Flame application

In this section, we provide a detailed overview of all functionalities available in the *F4F* application, explaining each tab individually.

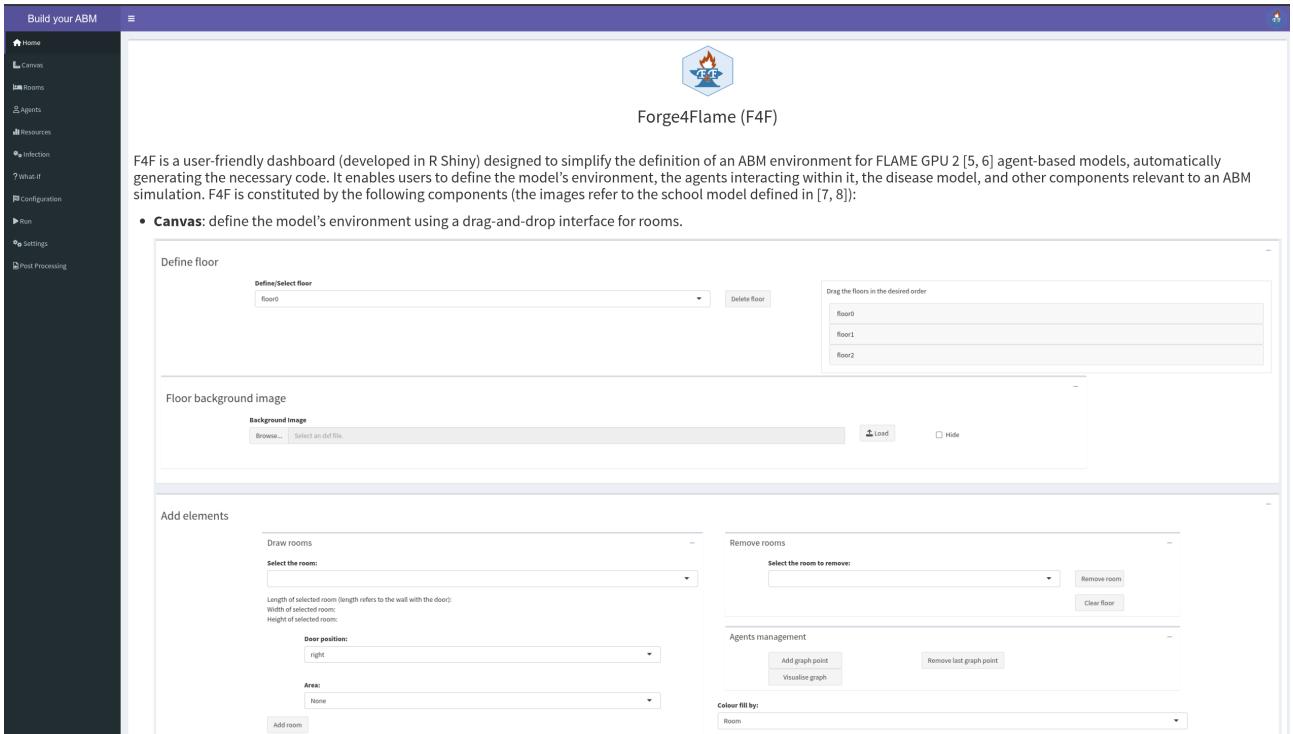


Figure 1: The **Homepage** page.

4.1 Homepage

Figure 1 illustrates the **Homepage** of *F4F*, providing a clear overview of the application. The left-hand menu organises *F4F*'s main functionalities, each designed to streamline various stages

of model creation.

4.2 Canvas

Figure 2 shows the initial **Canvas** page. In the panel **Define floor**, a drop-down menu can be used to select a previously defined floor or to create a new floor, giving it a name. The button *Delete floor* deletes the selected floor.

In the section *Draw rooms*, inside the panel **Add elements**, the user can add a previously defined type of room inside the canvas—by clicking on the button *Add room*—specifying the position of the door (none, bottom, left, top, right) and the area to which the room belongs (the user can define new areas or assign a room to no area). In the section *Remove rooms* the user can remove a room from the canvas by selecting it in the drop-down menu and clicking on the button *Remove room*. It is also possible to remove all the elements in the selected floor through the button *Clear floor*. The *Agent management* section is useful to create the graph of movements used by agents and to visualise it—see Section 6 for more details on the graph of movement. The *Add graph point* button allows the user to insert a new node into the agents' movement graph. It is particularly useful to insert these graph points near the corners of the "corridors". Additionally, each room's door and the room's centre area are represented by a node in the graph. In addition, it is possible to remove the last inserted graph point by clicking on the *Remove last graph point* button and to visualise all the graph's edges by clicking on the *Visualise graph* button. The drop-down menu named *Colour fill by* allows the user to colour rooms based on names, types, or areas. Finally, under this panel, there is the canvas, initially empty. Figure 3 shows the just-created canvas associated with floor *floor0*—with the panel **Add elements** closed. Each square in the canvas represents $1 \times 1 \text{ m}^2$.

Figure 4 shows the canvas of floor *floor0* with 16 rooms and 4 graph points. In particular, in this case, the colour of the room represents a previously defined room (that is, in the example light green represents the *room0*, dark blue the *room1*, gray the *fillingroom*, dark purple the *stairs*, and red the *spawnroom*). Moreover, Figure 5 shows the canvas of floor *floor1* with 15 rooms and 4 graph points (here we have two more defined rooms: dark pink representing the *room2*, and dark green the *room3*).

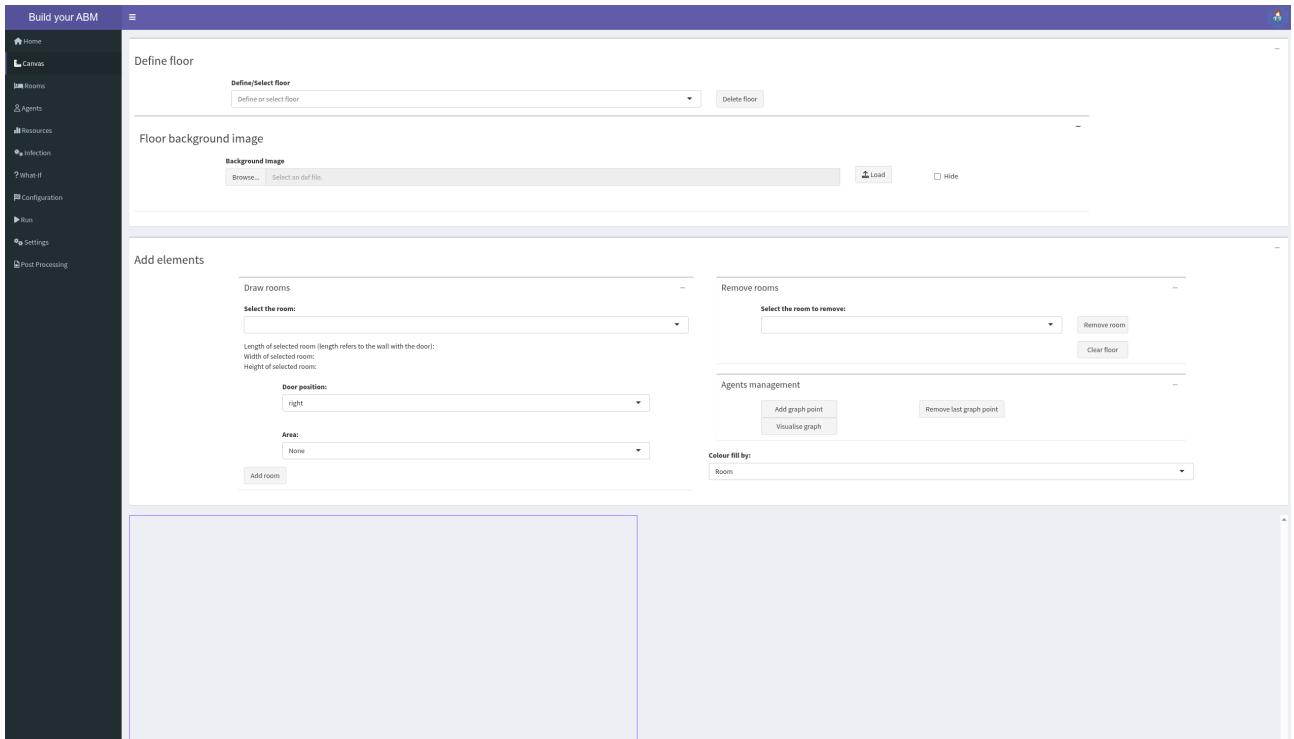


Figure 2: The **Canvas** page.

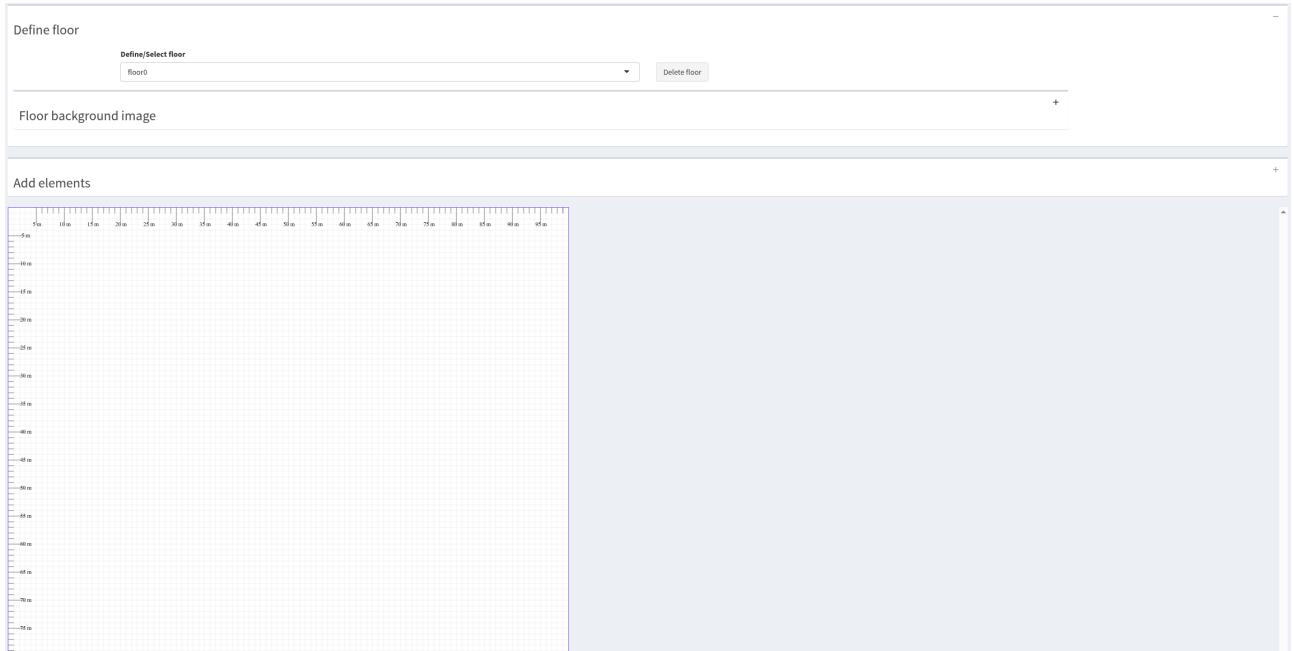


Figure 3: The **Canvas** page with the empty canvas of the *floor0* visualised.

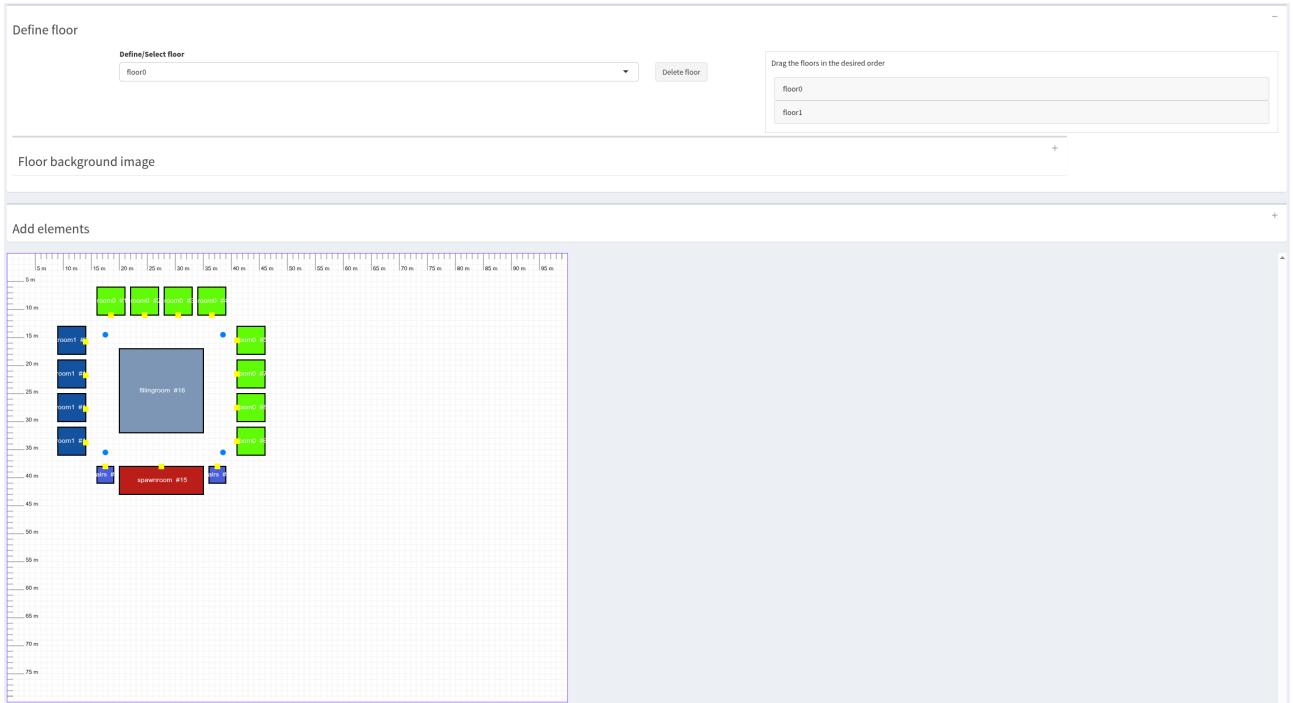


Figure 4: The **Canvas** page with some rooms and graph points inside the canvas *floor0*. Rooms are coloured by names.

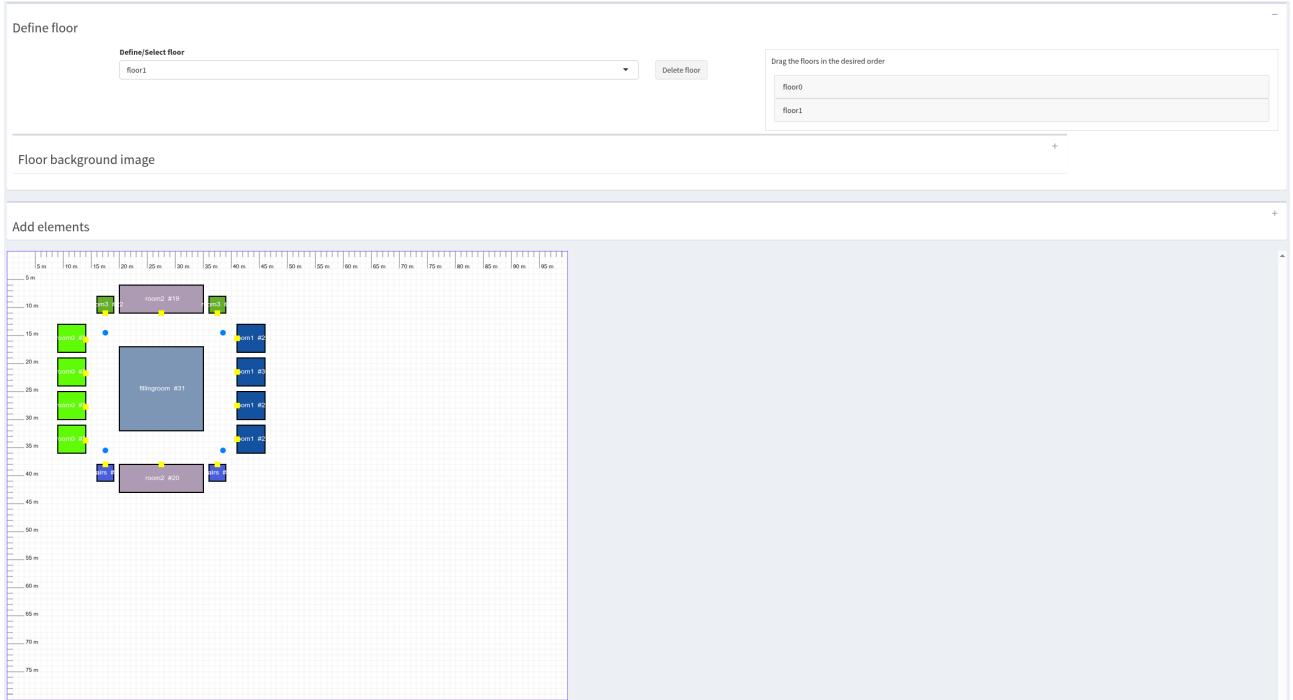


Figure 5: The **Canvas** page with some rooms and graph points inside the canvas *floor1*. Rooms are coloured by names.

4.2.1 Floor Plan Import as Background Image Support

F4F supports the import of architectural floor plans in the **DXF (Drawing Exchange Format)** file type as a background image. DXF is a widely used format for storing 2D and 3D

design data and is compatible with CAD software such as AutoCAD, QGIS, and other spatial or architectural design tools. This functionality aims to aid users in constructing realistic indoor environments by leveraging pre-existing building layouts.

Each modelled floor can be associated with a corresponding DXF. Upon loading a valid DXF file, the floor plan is rendered behind the canvas and serves as a visual reference during the manual design of the environment. F4F offers the capability to enable or disable the visibility of the background image interactively, facilitating a focus on the manual construction of the model after alignment is complete.

Each floor within the modelling environment can be associated with a distinct floor plan. Once a DXF is loaded:

- The floor plan is rendered as a background element on the design canvas.
- The rendering is always centred on the canvas to ensure consistent visual alignment.
- The canvas dimensions are automatically set according to the bounding box of the loaded plan, ensuring that all plan elements are fully visible.
- If multiple plans are loaded (for multiple floors), the canvas dimensions are updated to match the largest plan, guaranteeing consistency across floors.

The coordinate system used within the canvas adheres to a fixed spatial scale: **10 pixels correspond to 1 meter**. This convention allows for straightforward scaling of real-world measurements from DXF files to canvas units. Users are advised to ensure that their input files are scaled accordingly or to adjust dimensions prior to import.

To ensure compatibility and clarity, it is recommended that:

- **Each uploaded file should contain a single, clearly defined floor plan.** If multiple floors are to be used, each must be provided as a separate file.
- Only structural elements relevant to model construction (e.g., walls, openings) should be retained.
- Raster images used as background plans should be properly cropped and scaled to match the intended spatial resolution.

This capability significantly reduces manual errors and streamlines the setup process for simulation-ready spatial environments.

4.3 Rooms

The **Rooms** page allows the user to define the rooms that will be placed inside the canvas. There are two panels inside this page, as shown in Figure 6.

The first, **Define a new room**, allows creating the room from scratch, and the second, **Set colour legend**, allows modification of the colours associated with each defined room. To define a new room, the user must specify a *Name*—the name of a room is case insensitive, i.e. *room1* and *ROOM1* represent the same room—, the *Length*, the *Width*, and the *Height* of the room. Notice that the length refers to the wall with the door. The height of a room will not be visible in the canvas but will be displayed in the FLAME GPU 2 visualisation. It is essential for accurately implementing the aerosol contagion process—see Section 4.6 for more details.

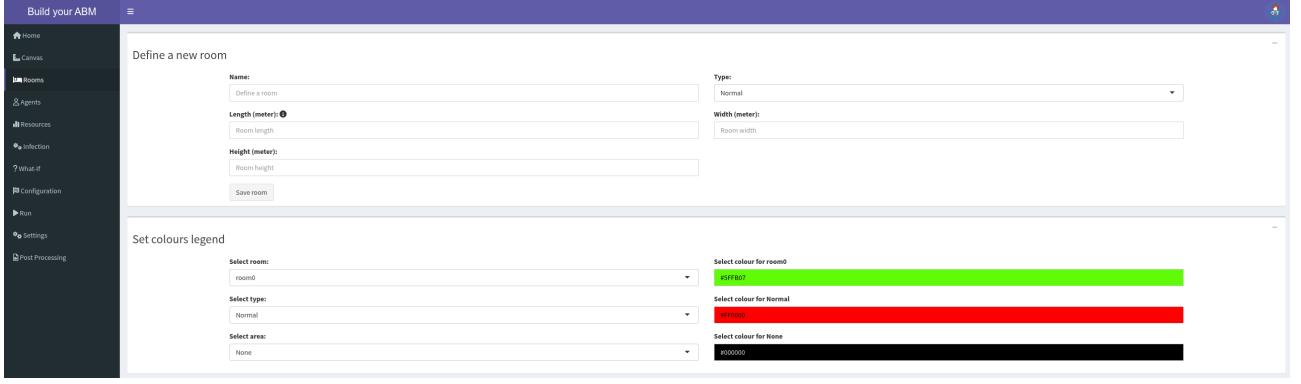


Figure 6: The **Rooms** page.

Rooms can be grouped into clusters based on a specific *Type*. This allows for the definition of a particular room type, instructing the agents to visit every room of that type. For instance, the user may define the generic type **bathroom**, defining a *small_bathroom* and a *large_bathroom*, instructing the agents to visit the generic **bathroom**. An example of the definition of the room *small_bathroom* can be seen in Figure 7. A set of default types is already defined inside our application:

- *Normal*: represents a generic room type.
- *Stairs*: room type used to move agents from one floor to another; stairs must be approximately in the same position on each floor.
- *Spawnroom*: room type from which agents enter and exit the environment; currently, only one Spawnroom can be placed inside the canvas.
- *Fillingroom*: room type with no door; can be used to simulate the walls and rooms of no interest.
- *Waitingroom*: room type used by agents as a waiting area before entering another room when no resources are available.

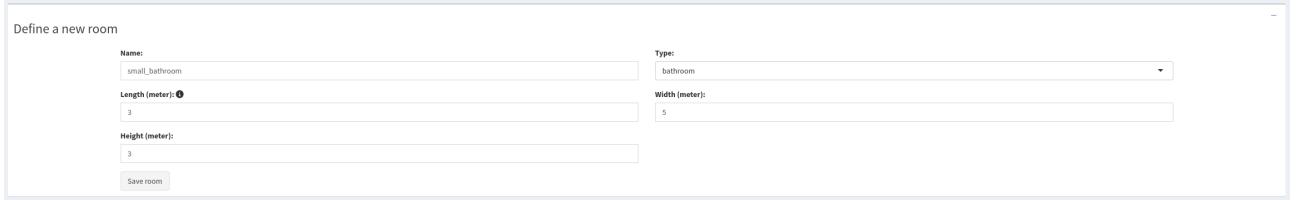


Figure 7: Definition of a room called *small_bathroom* associated to a type called **bathroom**.

The panel **Set colour legend** allows the user to define the colours of individual rooms, rooms of a specific type, or rooms in a particular area, based on the chosen option from the drop-down menu. The associated colour will be displayed next to the selected item. To change the colour, simply click on it. A colour map will appear, enabling the user to select their preferred colour, as shown in Figure 8. These colours will be the ones the user sees on the canvas and in the FLAME GPU 2 simulation.



Figure 8: Selection of a particular colour for the room *small_bathroom*.

4.4 Agents

Figure 9: The Agents page.

The **Agents** page serves to define the agents that will interact with the environment. It comprises several panels—see Figure 9:

- In the **Agent definition** panel, the user can define all possible agent types. Each agent is defined by a name (*Agent name*). For each specific agent, the user can define details like determined and random flows and entry times, which will be covered in the following sections. The user can also copy all information concerning an agent using the drop-down menu *Copy information from*, which allows the user to select an existing agent and then modify it as needed. Additionally, the user can specify the number of agents to include in the model by entering the desired amount in *Number of agents*. An example of this process is shown in Figure 10.

Figure 10: An example of the definition of an agent named *Neurologist*, with a total of 10 agents of that type in the model.

- In the Determined flow panel, the user can define one or more flows that an agent must follow within the model. There are three key elements to configure:

1. In the *Type* drop-down menu, the user can select the room to add to the flow. These rooms are those that have been added to the canvas—by type and area.
2. In the *Activity* drop-down menu, the user can select the intensity of the activity that the agent performs in the selected room. This setting is crucial for calculating the virus concentration within a room: higher activity intensity leads to greater virus spread.
3. The third and last element defines the duration of the task performed by the agent. The user can specify a deterministic time, which represents the exact number of minutes the agent spends on the task, or define the duration using a stochastic distribution—i.e., Exponential, Uniform, or Truncated Positive Normal. The distribution type can be selected from a drop-down menu, and its parameters can be entered below.

Rooms are added to the flow using the *Add Room* button, and the entire flow will be displayed below. The first room in the flow will appear at the top, with the agent following the subsequent rooms in descending order. Both the first and last rooms in the flow **must** be the room with type *Spawnroom* and area *None* to allow the agent to enter and exit the environment—the time associated with the last *Spawnroom* is **negligible**, it will be automatically computed based on the next entry time. The order of rooms can be adjusted by selecting and moving them up or down within the flow. Additionally, the last room can be removed using the *Remove last room* button. Multiple flows can be defined for a single type of agent—except for agents with a daily rate, which can have only one flow, as explained below. An example of a defined flow is shown in Figure 11.

The screenshot shows the 'Determined flow' panel with the following interface elements:

- Top Controls:** Type: dropdown, Activity: dropdown, Deterministic/Stochastic radio buttons, Fixed deterministic value: input field (value), Add room, Remove last room.
- Middle Controls:** Add flow, Remove flow, 1 flow.
- Flow List:** A scrollable list titled 'Drag the rooms in the desired order' containing the following items:
 - Spawnroom-None - Deterministic 10 min - Light
 - Normal-None - Exponential 100 min - Quite Hard
 - Normal-Area1 - Uniform a = 50; b = 250 min - Hard
 - Normal-Area0 - Truncated Positive Normal Mean = 100; Sd = 10 min - Hard
 - (2) Spawnroom-None - Deterministic 10 min - Light

Figure 11: An example of a determined flow. Imagine that the agent is scheduled to enter the environment at 10:00 A.M. Then, it will *i*) remain outside the environment until 10:10 A.M., while engaging in a light activity, *ii*) move to a room of type *Normal* with area *None*, where it will perform a quite hard activity for a duration determined by an exponential distribution with an average of 100 minutes, *iii*) proceed to a room of type *Normal* with area *Area1*, where it will perform a hard activity for a duration determined by a uniform distribution with $a = 50$ minutes and $b = 250$ minutes, *iv*) move to a room of type *Normal* with area *Area0*, where it will perform a hard activity for a duration determined by a truncated positive normal distribution with a mean of 100 minutes and standard deviation of 10 minutes, and finally *v*) leave the environment.

- In the Random flow panel, the user can add random events that interrupt the agent's determined flow, i.e. directing the agent to an additional room—a random event should occur infrequently and last for a short duration (a few minutes). The elements are similar to those in the previous section, including the *Type* drop-down menu for selecting rooms—by type and area—, the *Activity* drop-down menu for selecting the activity type, and the definition of the time spent in that room. The *Weight* field allows the user to specify the probability of the event occurring by entering a number between 0 and 1. Users can specify the time window for the event by setting the *Initial time* and *Ending time* fields. By default, the event spans the entire day (00:00 to 23:59). Alternatively, users can define a specific time range—for example, an event like lunch occurring only during lunchtime. By default, the agent will follow the determined flow without any events. When a new event (e.g., going to the toilet) is added with a probability p , the probability of following the determined flow in each simulation step becomes $1 - p$. A random event should happen rarely and last only a few minutes. If an event lasts too long, it may be cut short due to the completion of a block in the determined flow. For example, consider an event e with a duration of t_e , while the agent is required to remain in its current position for t_{df} —where df represents the current block in the determined flow. If $t_e > t_{df}$, the agent will experience the event for t_{df} to maintain consistency with the determined flow. In each simulation step, at most one event can occur, and the probabilities are to be considered **per simulation step**—for more details on the simulation step, see Section 4.8. An agent is restricted from triggering new events while actively handling an ongoing event to avoid excessively long or infinite event loops. New events can only be initiated once the current event has been completed. In addition, an event can be removed by simply clicking on it—an alert will appear asking for confirmation to delete or cancel the action. An example of a random flow is shown in Figure 12.

Room	Distribution	Activity	Time	Weight	Time Slot
Do nothing	Deterministic	Light	0	0.9992	00:00 - 23:59
Normal-Area2	Deterministic	Light	15	8e-04	00:00 - 23:59

Figure 12: An example of random flow. At each simulation step, the agent has a probability of 0.0008 to move to a room of type *Normal* with area *Area2* with light activity and will spend exactly 15 minutes there. The event can occur in the specified time slot (in this case, 00:00 to 23:59).

- The last panel is called **Entry flow**. It allows the user to define when the agent enters the environment and starts its determined flow. There are two ways to define an agent's entry flow:
 - By a **daily rate**: the user can define an entrance rate for the agent using either a deterministic value or a stochastic distribution, along with specifying an initial and a final generation time. Agents will be generated inside this time window. If a deterministic value is used, the exact number of agents specified will be generated; otherwise, the number of agents created will be determined according to the stochastic distribution parameters. Generated agents will follow the determined flow and

leave the environment upon completion—these agents can have only one determined flow. Additionally, the user can set various entry rates for different time windows and specify particular days of the week for agent generation. See Figures 13 and 14 for an example of a definition of a daily rate entry flow.

- By a **time window**: in this case, the number of agents is specified in the **Agent definition** panel—see Figure 10. The user sets the entry time and the specific days of the week when the agent will enter the environment—it is possible to define multiple time slots. Additionally, the user can define different flows—see Figure 11—associating one of them with each time slot. These agents will repeat the selected flow each week throughout the simulation. The total duration of a determined flow can only be estimated (especially if the user employs stochastic distributions). We check for possible overlaps among the entry flows using the average of the selected distributions. In some cases, overlaps may occur. In these cases, when the agent finishes the current determined flow, it will start the next one with some delay. Our suggestion is to avoid placing different entry flows too close to each other. See Figure 15 for an example of the definition of a time window entry flow.

The screenshot shows the 'Entry flow' configuration interface. Under 'Select type of entrance:', the 'Daily Rate' radio button is selected. The 'Entrance rate:' section includes 'Deterministic' and 'Stochastic' buttons, with 'Deterministic' currently active. The 'Distribution:' dropdown is set to 'Uniform'. The 'a:' field contains '10' and the 'b:' field contains '50'. The 'Initial generation time:' field shows '10:30' and the 'Final generation time:' field shows '19:45'. On the right, a 'Select Days of the Week' section lists Monday through Friday with checkboxes checked. At the bottom are 'Add slot' and 'Remove slot' buttons.

Figure 13: An example of a daily rate entry flow (slot 1). Agents will be generated between 10:30 A.M. and 7:45 P.M. on every weekday except Tuesday and Sunday, with their number determined stochastically using a uniform distribution with $a = 10$ agents and $b = 50$ agents.

The screenshot shows the 'Entry flow' configuration interface. Under 'Select type of entrance:', the 'Daily Rate' radio button is selected. The 'Entrance rate:' section includes 'Deterministic' and 'Stochastic' buttons, with 'Deterministic' currently active. The 'Fixed deterministic value:' field contains '5'. The 'Entry time:' field shows '12:00' and the 'Exit time:' field shows '22:00'. On the right, a 'Select Days of the Week' section lists Tuesday and Sunday with checkboxes checked. At the bottom are 'Add slot' and 'Remove slot' buttons.

Figure 14: An example of a daily rate entry flow (slot 2). Five agents will be generated between 12:00 A.M. and 10:00 P.M. on Tuesday and Sunday.

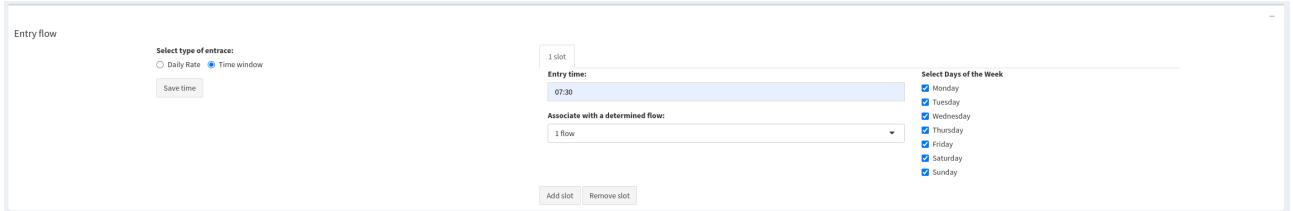


Figure 15: An example of a time window entry flow. Agents will enter the environment at 7:30 A.M. every day of the week, following the first flow.

4.5 Resources

The **Resources** page allows the user to specify the maximum number of agents that can occupy a room simultaneously—see Figure 16. It also provides options to set limits based on the type of agent.

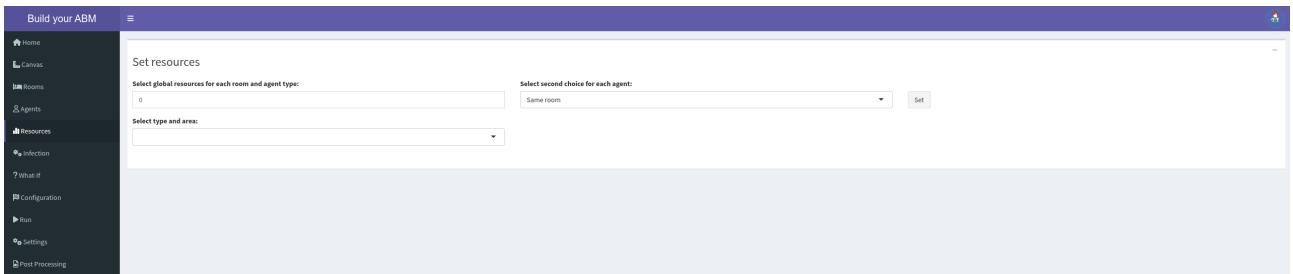


Figure 16: The **Resources** page.

This page consists solely of the **Set Resources** panel, as illustrated in Figure 17. Within this panel, two global values can be configured:

- **Maximum number of agents:** users can specify globally the maximum number of agents allowed to enter each room in the environment. This is particularly useful when the user is not interested in modelling resource limitations and prefers to set them to high values.
- **Global alternative for room entry:** if an agent is unable to enter a room in its flow due to capacity constraints, the user can define an alternative option based on the associated resources. The available alternatives include:
 - *Same room:* search for another room of the same type and area.
 - *Skip room:* proceed to the next room in the flow or skip the event (in the case an event occurred).
 - *Another room:* select a different room with a different type and area.

To set this value, simply enter the desired number in the input field and press the *Set* button.

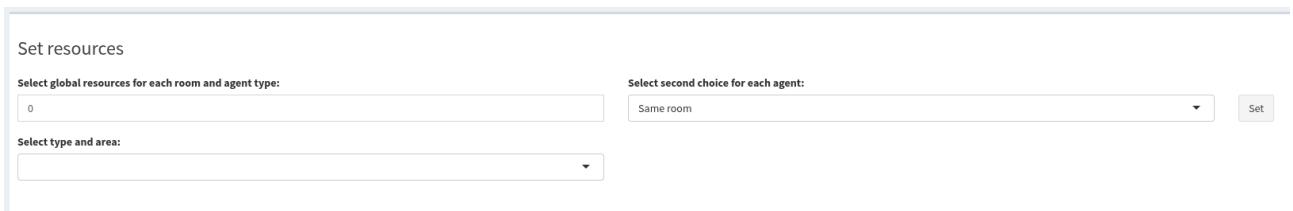


Figure 17: The *Set resources* panel.

The third element of the section is a drop-down menu labelled *Select Type and Area*, which allows the user to choose a room to associate resources with, categorised by type and area. By selecting a room based on its type and area, the user can modify the associated resources. As shown in Figure 18, a table displays all rooms of the selected type and area that are associated with at least one agent's flow.

The screenshot shows a user interface titled 'Set resources'. At the top left is a dropdown for 'Select global resources for each room and agent type' with the value '1000'. To its right is another dropdown for 'Select second choice for each agent' with the value 'Same room', and a 'Set' button. Below these is a dropdown for 'Select type and area' set to 'Normal-Area1'. A table follows, with columns for 'Room', 'Maximum' (with values 'room1: 1000' and 'room2: 1000'), 'Neurologist' (with values '1000' for both rooms), and 'Surgeon' (with values '1000' for both rooms). Below the table are two dropdowns: 'Select second choice room in Determined Flow for Neurologist' and 'Select second choice room in Random Flow for Surgeon', both currently set to 'Same room'.

Figure 18: The resources associated with rooms of type *Normal* and area *Area1* are displayed. This room appears in the predefined flows of the *Neurologist* agent and in the random flow of the *Surgeon* agent.

For each room, the user can view and modify the number of agents allowed to enter by double-clicking the respective field. The total number of agents permitted in the room can be adjusted in the *Maximum* column. Additionally, the user can specify the allowed number of agents for each type, with each type displayed in a separate column corresponding to the room. If an agent finds the room full when attempting to enter, the user can specify an alternative room using the *Select second choice room in Determined/Random Flow for Agent* drop-down menu for each agent. The user can assign different alternative rooms for an agent's determined and random flow, as shown in Figure 18.

4.6 Infection

The **Infection** page—see Figure 19—allows the user to choose from eight different disease compartmental models:

- SIR: Susceptible-Infected-Recovered.
- SIRD: Susceptible-Infected-Recovered-Deceased.
- SEIR: Susceptible-Exposed-Infected-Recovered.
- SEIRD: Susceptible-Exposed-Infected-Recovered-Deceased.
- SIRS: Susceptible-Infected-Recovered-Susceptible.
- SIRDS: Susceptible-Infected-Recovered-Deceased-Susceptible.
- SEIRS: Susceptible-Exposed-Infected-Recovered-Susceptible.
- SEIRDS: Susceptible-Exposed-Infected-Recovered-Deceased-Susceptible.

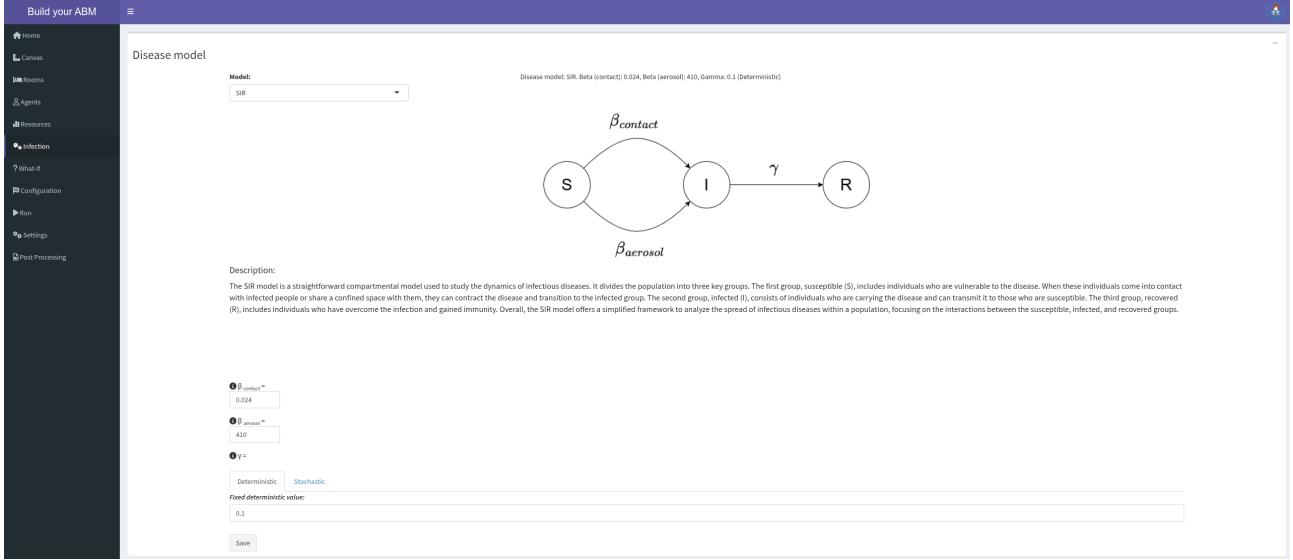


Figure 19: The **Infection** page with the SIR compartmental model shown.

A brief description and a graphical schema are provided for each disease model. Based on the chosen model, the user must specify values for the various parameters of that model. In particular, all the disease models share three parameters: *i*) $\beta_{contact}$, which represents the contamination risk associated with infection from close-range contacts, as modelled in [14]—for example, for the COVID-19 disease the value to use is 0.024—, *ii*) $\beta_{aerosol}$, which represents the risk constant associated with infection due to aerosol, as modelled in [12]—for example, for the COVID-19 disease the value to use is 410—, and *iii*) γ , which represents the recovery rate. **Note that $\beta_{contact}$ and $\beta_{aerosol}$ are not the infection rates.** For each parameter, except $\beta_{contact}$ and $\beta_{aerosol}$, the user can choose a deterministic value or a stochastic distribution. After specifying the parameters, click on the *Save* button to save them.

In Figure 20, the SEIRDS model is shown. This case is characterized by six parameters: *i*) $\beta_{contact}$, *ii*) $\beta_{aerosol}$, *iii*) γ , *iv*) α , which represents the incubation rate, *v*) λ , which represents the fatality rate, and *vi*) ν , which represents the end-of-immunization rate.

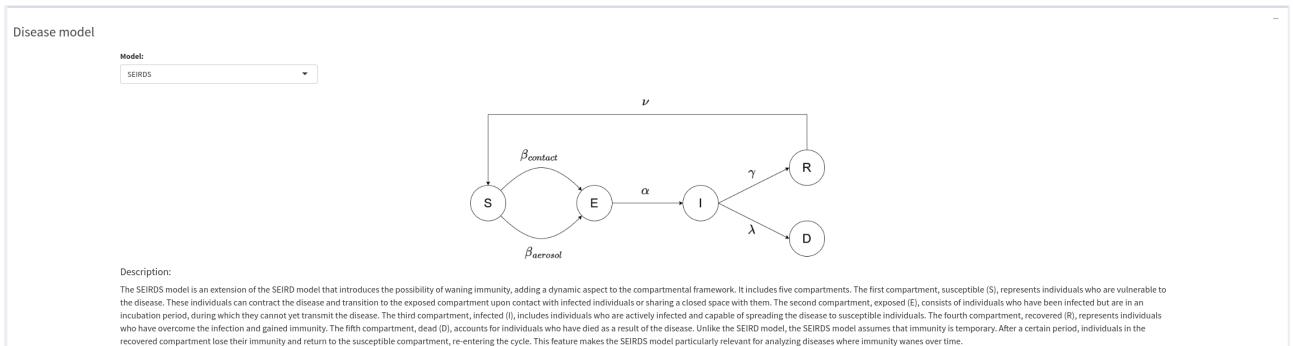


Figure 20: The SEIRDS compartmental model.

4.7 What-If

The screenshot shows the 'Countermeasures' section with several tables for different measures. The first table for 'Measure' has columns for 'Measure' (All), 'Type' (All), 'Parameters' (All), 'From' (All), and 'To' (All). The second table is identical. Below these are sections for 'Ventilation', 'Masks', 'Vaccination', 'Swabs', 'Quarantine', and 'External screening'. The 'Virus parameters' section contains a table for 'Virus' with 'Virus variant factor' set to 1 and 'Virus severity' set to 0.22. The 'Initial infected' section has two input fields: 'Initial infected' (with a dropdown menu) and 'Initial infected' (with a numerical input field set to 1).

Figure 21: The **What-If** page.

The **What-If** page—see Figure 21—enables the user to conduct various what-if scenarios considering different *i*) levels of environmental ventilation, *ii*) masks policies, *iii*) levels of vaccination coverage and efficacy, *iv*) screening policies, *v*) quarantine policies, *vi*) external screening policies, *vii*) virus variants and severity of symptoms, *viii*) initial infected types and numbers, and *ix*) levels of outside contagion. The page is divided into two macro-sections: **Countermeasures** and **Virus Parameters**. If the user does not specify a particular countermeasure, it will not be applied. For the virus parameters, the default values will be used—see below.

For each countermeasure, the user must specify the period during which it will be in effect. For instance, the user can specify a particular type of mask with a defined fraction of agents wearing it from day d_1 to d_2 and apply a different setting from $d_2 + 1$ to d_3 —with $d_1 < d_2 < d_3$. Clicking the *Save* button the specified settings will be saved—this applies to both the countermeasures and the virus parameters. In details the **Countermeasures** section contains:

- The **Ventilation** panel allows the user to define either a global ventilation level—Figure 22—, which is the same for each room, or specific ventilation levels—Figure 23— which can differ for each room [12]. There are 7 possible ventilation levels in ACH¹, i.e. 0 (no ventilation), 0.3 (poorly ventilated), 1 (domestic), 3 (offices/schools), 5 (well-ventilated), 10 (typical maximum), and 20 (hospital setting). The ventilation level has an impact on the aerosol contagion process—see Section 4.6.

¹For instance, 3 Air Changes per Hour (ACH) means that in 1 hour 300.000 L (or analogous 300 squared meters) of external air are entered into the considered room.

Figure 22: Setting ventilation globally from day 1 to 10.

Figure 23: Setting ventilation specifically to rooms with type *Normal* and area *Area1* from day 1 to 10.

- The **Masks** panel allows the user to define either a global mask policy—Figure 24—which is the same for each agent, or a specific mask policy—Figure 25—which can differ for each agent. There are 3 different types of masks [20], i.e. no mask (0% of exhalation and inhalation efficacy), surgical mask (59% of exhalation and inhalation efficacy) and FFP2 mask (90% of exhalation and inhalation efficacy). The user can also specify the fraction of agents that will wear the selected mask, either globally or for a specific agent.

Figure 24: Setting mask policy globally from day 1 to 10.

Figure 25: Setting mask policy specifically to agents with type *Surgeon* from day 1 to 10.

- The **Vaccination** panel allows the user to define either a global vaccination policy—Figure 26—which is the same for each agent, or a specific vaccination policy—Figure 27—which can differ for each agent. The user must specify the fraction of agents that will be vaccinated and the efficacy of the vaccines. We define vaccination efficacy as the proportion of vaccinated agents that are effectively immune to infection, while the remainder remain susceptible despite being vaccinated. It is also necessary to set a vaccine coverage, which determines the number of days the vaccination will protect the agent. The user can configure this using either a deterministic value or a stochastic distribution, as described previously. In this case, the countermeasure will be applied on the specified day. For instance, if the user specifies a fraction f of vaccinated agents with an efficacy e and a coverage c on day d , those agents will be vaccinated on that day. In this case, there is no *From* and *To* period.

Vaccination

Vaccination:
 Global
 Different for each agent

Fraction of vaccinated agents: Vaccine efficacy:

Vaccine coverage (day):
 Deterministic Stochastic
 Distribution:
 Value:
 At (day):

Figure 26: Setting vaccination policy globally at day 5.

Vaccination

Vaccination:
 Global
 Different for each agent

Agent: Fraction of vaccinated agents: Vaccine efficacy:

Vaccine coverage (day):
 Deterministic Stochastic
 Distribution:
 Value:
 At (day):

Figure 27: Setting vaccination policy specifically to agents with type *Surgeon* at day 5.

- The **Swabs** panel allows the user to define either a global policy—Figure 28—, which is the same for each agent, or a specific policy—Figure 29—which can differ for each agent. The user must specify the sensitivity and specificity of the swabs, as well as the frequency (deterministic or stochastic) at which swabs are administered to each agent.

Swabs

Swab:
 Global
 Different for each agent

Sensitivity: Specificity:

A swab every how many days?
 Deterministic Stochastic
 Fixed deterministic value:

From (day): To (day):

Figure 28: Setting swab policy globally from day 3 to 10.

Swabs

Swab:
 Global
 Different for each agent

Sensitivity: Specificity:

Agent: Swab:
 No swab
 Swab

A swab every how many days?
 Deterministic Stochastic
 Fixed deterministic value:

From (day): To (day):

Figure 29: Setting swab policy specifically to agents with type *Neurologist* from day 3 to 10.

- The Quarantine panel allows the user to define either a global policy—Figure 30—, the same for each agent, or a specific policy—Figure 31—which can differ for each agent. The user must specify the duration of the quarantine in days (deterministic or stochastic) and the room used to put in quarantine agents with severe cases—the default is outside the environment. The user can also choose a swab policy, enabled only during quarantine, by specifying the frequency (deterministic or stochastic) at which swabs are administered to each agent, the sensitivity and the specificity.

This screenshot shows the Quarantine panel with global settings applied to all agents. The 'Quarantine' section is set to 'Global'. The 'Quarantine days' section uses a 'Deterministic' fixed value of 10. The 'Quarantine room for severe cases' is set to 'Spawnroom-None'. The 'Swab' section is enabled ('Swab' radio button selected), with sensitivity and specificity both set to 1. The 'A swab every how many days?' section is set to a deterministic value of 2. The 'From (day)' and 'To (day)' fields are set from 1 to 10. A 'Save' button is visible in the bottom right corner.

Figure 30: Setting quarantine policy globally from day 1 to 10.

This screenshot shows the Quarantine panel with specific settings for the 'Surgeon' agent. The 'Agent' dropdown is set to 'Surgeon'. The 'Quarantine' section is set to 'Different for each agent'. The 'Quarantine days' section uses a 'Deterministic' fixed value of 7. The 'Quarantine room for severe cases' is set to 'Normal-Area2'. The 'Swab' section is enabled ('Swab' radio button selected), with sensitivity and specificity both set to 1. The 'A swab every how many days?' section is set to a deterministic value of 2. The 'From (day)' and 'To (day)' fields are set from 1 to 10. A 'Save' button is visible in the bottom right corner.

Figure 31: Setting quarantine policy specifically for agents with type *Surgeon* from day 1 to 10.

- The External screening panel allows the user to define either a global policy—Figure 32—the same for each agent, or a specific policy—Figure 33—which can differ for each agent. In particular, two different probabilities of external screening can be applied to agents each day: the probability of testing an agent due to activities that involve screening campaigns (e.g., sports practice) and the probability of testing an agent outside the environment due to the presence of symptoms.

This screenshot shows the External screening panel with global settings applied to all agents. The 'External screening' section is set to 'Global'. The 'Screening campaigns' field has a value of 0.0075. The 'Symptoms' field has a value of 0.03. The 'From (day)' and 'To (day)' fields are set from 1 to 10. A 'Save' button is visible in the bottom right corner.

Figure 32: Setting external screening policy globally from day 1 to 10.

External screening

External screening:
 Global
 Different for each agent

Agent: Neurologist Screening campaigns: 0 Symptoms: 0

From (day): 1 To (day): 10

Figure 33: Setting external screening policy specifically for agents with type *Neurologist* from day 1 to 10.

At the beginning of the **What-If** page, all countermeasures set by the user are displayed, as shown in Figure 34, in the *Saved Countermeasures* panel. The first table contains countermeasures related to the rooms (i.e., ventilation), while the second table contains countermeasures related to agents. To remove one of the countermeasures, the user must click on the relevant entry, and an alert will appear asking for confirmation to delete or cancel the action. It is also possible to display only specific countermeasures using the filter associated with each column. An example is shown in Figure 35, where only the mask policy is displayed. **Specific countermeasures will take precedence over global ones.** For instance, as shown in Figure 34, 50% of agents will wear a surgical mask from day 1 to day 10, with the exception of surgeons, 70% of whom will wear an FFP2 mask.

Measure Type Parameters From To

All	All	All	All	All
Ventilation	Global	1	1	10
Ventilation	Normal-Area	3	1	10

Measure Type Parameters From To

All	All	All	All	All
Mask	Global	Type: Surgical mask; Fraction: 0.5	1	10
Mask	Surgeon	Type: FFP2 mask; Fraction: 0.7	1	10
Vaccination	Global	Efficacy: 0.3; Fraction: 0.5; Coverage Dist.Days: Deterministic, 60, 0	5	5
Vaccination	Surgeon	Efficacy: 1; Fraction: 0.7; Coverage Dist.Days: Deterministic, 60, 0	5	5
Swab	Global	Sensitivity: 1; Specificity: 1; Dist: Deterministic, 5, 0	3	10
Swab	Neurologist	Sensitivity: 1; Specificity: 1; Dist: Deterministic, 2, 0	3	10
Quarantine	Global	Dist.Days: Deterministic, 10, 0; Q.Room: Spawndroom-None; Sensitivity: 1; Specificity: 1; Dist: Deterministic, 2, 0	1	10
Quarantine	Surgeon	Dist.Days: Deterministic, 7, 0; Q.Room: Normal-Area; Sensitivity: 1; Specificity: 1; Dist: Deterministic, 2, 0	1	10
External screening	Global	First: 0.0075; Second: 0.03	1	10
External screening	Neurologist	First: 0; Second: 0	1	10

Figure 34: Saved countermeasures listed in the panel *Countermeasures*.

Search:

Measure Type Parameters From To

All	All	All	All	All
Ventilation	Global	1	1	10
Ventilation	Normal-Area	3	1	10

Search:

Measure Type Parameters From To

["Mask"]	All	All	All	All
Mask	Global	Type: Surgical mask; Fraction: 0.5	1	10
Mask	Surgeon	Type: FFP2 mask; Fraction: 0.7	1	10

Figure 35: Example of filter on the *Measure* column.

On the other hand, the **Virus parameters** section contains:

- The **Virus** panel allows the user to define the virus—or the variant of the virus—to model and the severity of symptoms—Figure 36. Use the value 1 as *Virus variant factor* if there are no variants or if the user wants to model the base variant—in [20] the user can find some examples for COVID-19. The *Virus severity* represents the probability of showing severe symptoms—in [23] the user can find an example of COVID-19.

Virus

Virus variant factor: 1

Virus severity: 0.22

Save

Figure 36: Definition of virus parameters.

- The **Initial infected** panel allows the user to define the number of agents that will be initially infected. There are three options:
 1. Random: if the number of initial infected is n , then these n agents will be randomly selected from all existing agents—Figure 37.
 2. Global: if the number of initial infected is n , then n agents of each agent type will be initially infected—Figure 38.
 3. Specific: for each agent type, the user can indicate the number of agents of that type that will be initially infected—Figure 39.

Initial infected

Initial infected: 2

Random

Save

Figure 37: Setting random initial infected.

Initial infected

Initial infected: 5

Global

Save

Figure 38: Setting initially infected globally.

Initial infected

Initial infected: 1

Different for each agent

Agent: Neurologist

Save

Figure 39: Setting initially infected specifically.

- The **Outside contagion** panel—Figure 40—allows users to model infections occurring outside the environment as a stochastic arrival process. This setup aligns the rate of external infections with the prevalence of the disease in the broader population. In particular, a probability is applied to each susceptible agent every day, representing the possibility of infection outside the environment. If the user wishes to model the infection outside the environment, they must upload a CSV file containing two columns: *day*, representing the simulation day, and *percentage_infected*, which indicates either the actual number of infected individuals in the considered population or the probability of acquiring the infection outside the environment on that day. If the former is provided, the user must specify the total population in the *Population* field; otherwise, they should set the population to 1. A plot is displayed after the file is uploaded—Figure 41.

The screenshot shows a user interface for 'Outside contagion'. At the top left, there's a label 'Outside contagion' with a help icon. Below it is a file upload section with a 'Browse...' button, a text input field 'Select an csv file.', and a 'Load' button. Underneath is a 'Population:' label with a text input field containing the value '1'.

Figure 40: Outside contagion panel.

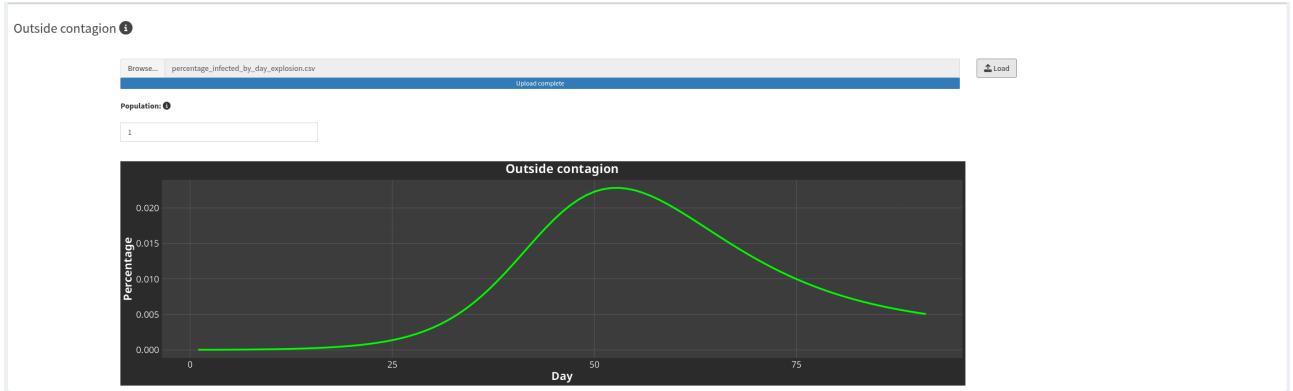


Figure 41: Example of a plot generated after uploading an outside contagion file.

At the beginning of the *Virus Parameters* section, all configured virus parameters are displayed in the *Saved Virus Parameters* panel, as shown in Figure 42. The first table shows the virus variant factor and virus severity, while the second one displays the initial infection set by the user. Entries can only be removed from the second table by clicking on them. An alert will appear, prompting the user to confirm the deletion or cancel the action. **Specific initial infections will take precedence over global ones, while randomly assigned infections are considered separately.**

The screenshot shows the 'Saved Virus Parameters' panel. It contains two tables. The first table has columns 'Variant', 'Severity', and 'Number'. It shows one entry: Variant 1, Severity 0.22, and Number 2. The second table has columns 'Type' and 'Number'. It shows three entries: Type Random (Number 2), Type Global (Number 5), and Type Neurologist (Number 1).

Saved Virus Parameters				
Variant	Severity	Type	Number	
1	0.22	Random	2	
		Global	5	
		Neurologist	1	

Figure 42: Saved countermeasures listed in the panel *Virus parameters*

4.8 Configuration

The screenshot shows the 'Configuration' page. On the left is a sidebar with navigation links: Home, Canvas, Rooms, Agents, Resources, Infection, What-if, Configuration, Run, Settings, and Post Processing. The main area is titled 'Configuration parameters' and includes fields for 'Simulation days' (set to 10), 'Initial day' (radio buttons for Monday through Sunday, with Monday selected), 'Initial time' (set to 07:00), and 'Step' (set to 60).

Figure 43: The Configuration page.

On this page—see Figure 43—the user can configure various settings for the simulation:

- *Simulation days*: define the duration of the simulation in days.
- *Initial day*: set the start day of the week for the simulation.
- *Initial time*: determine the start time of the day for the simulation.
- *Step*: define the duration of a step in seconds; possible values are dividers of 60 seconds.

4.9 Run

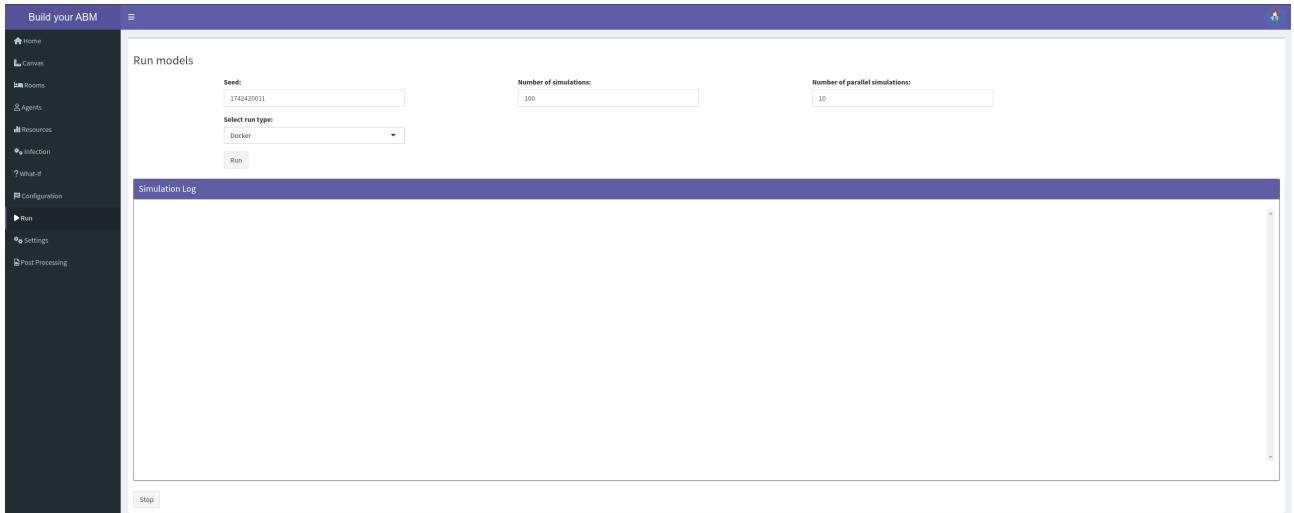


Figure 44: The **Run** page.

On this page—see Figure 44—the user can configure various settings for the simulation:

- *Seed*: specify the starting seed for the simulation.
- *Number of simulations*: specify the number of simulations to run with a specific configuration of parameters.
- *Number of parallel simulations*: specify the number of simulations to run in parallel to speed up the computation.

Moreover, as shown in Figure 45, the user can choose how to run the simulation they have just created using FLAME GPU 2. There are three options:

- Launch the simulation using Docker (see Section 2.4); in this case, visualisation will not be available.
- Launch the simulation locally with visualisation.
- Launch the simulation locally without visualisation.

These three options are available only when running *F4F* without Docker. If *F4F* is run using Docker Compose, the only available option will be *Docker*, while if *F4F* is run using Docker (without Compose), it will not be possible to run simulations using the *Run* button.

When the *Run* button is clicked, the system performs several checks, and the user must specify a location to save the results by clicking the *Select Folder* button, as shown in Figure 46—when running the process with Docker Compose, the results are automatically saved

in the directory containing the `docker_compose.yml` file. Additionally, the user must enter a name for the directory where the results will be stored. Finally, the user can terminate the running process by clicking the *Stop* button.

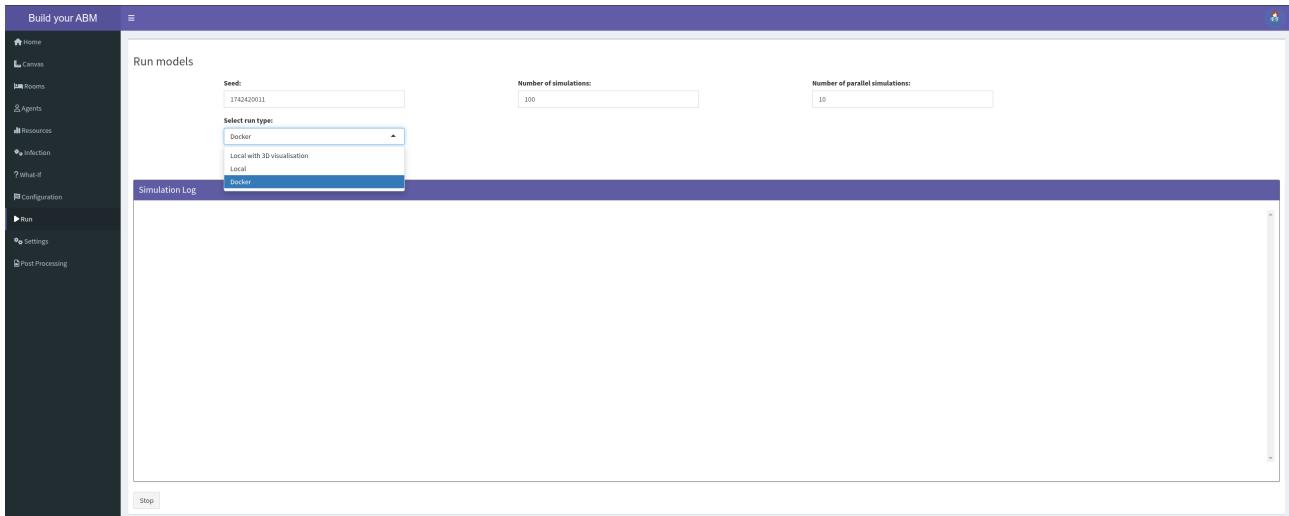


Figure 45: Section where the user can select how to run their simulation.

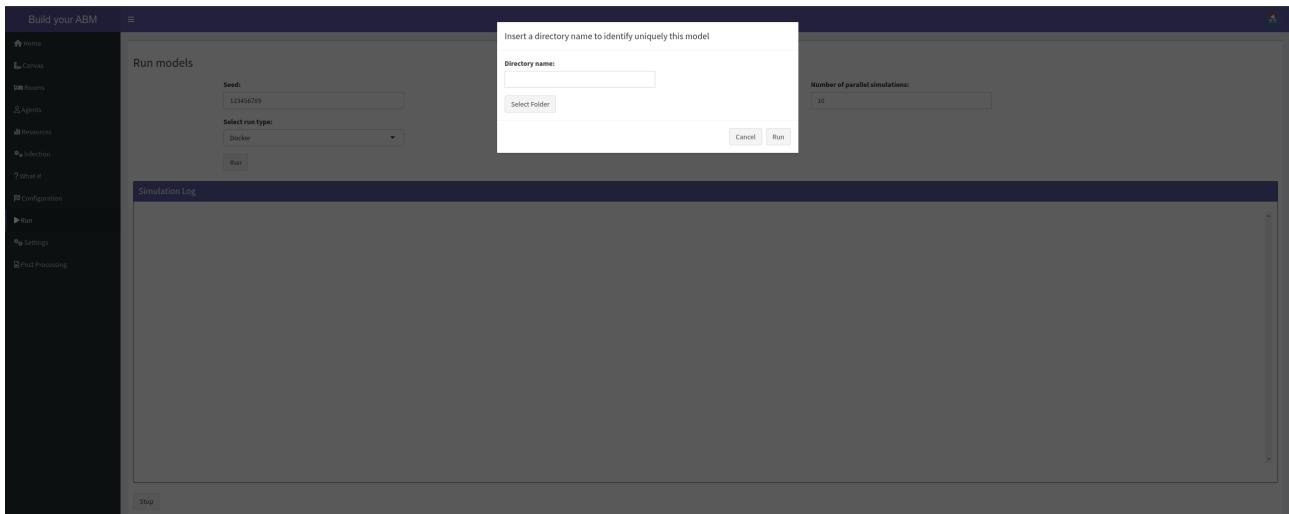


Figure 46: Pop-up displayed when the *Run* button is clicked.

4.10 Settings

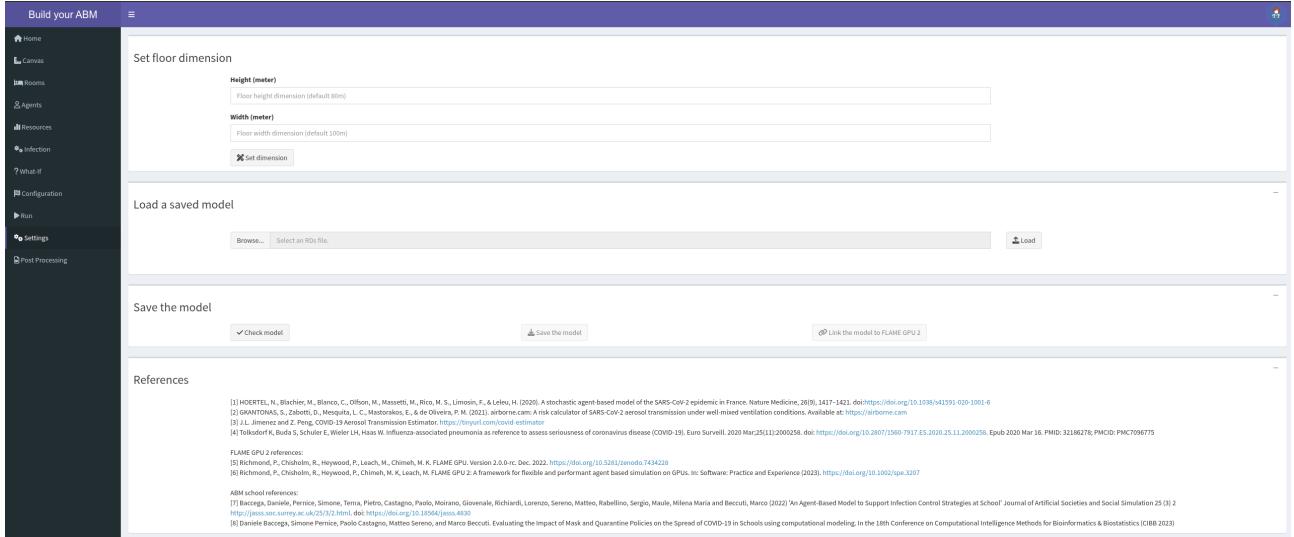


Figure 47: The **Settings** page.

Inside the **Settings** page—see Figure 47—the user can find three different panels:

- In the **Set floor dimension** panel, the user can adjust the dimensions of the canvas displayed on the **Canvas** page. The dimensions are specified in meters. The default dimensions are 80 meters in height and 100 meters in width. To set the desired dimensions, simply enter the values and press the *Set dimension* button.
- In the **Load a saved model** panel, the user can load a previously defined model. Click the *Browser* button to navigate through files and select the desired file. Once the correct file is chosen, click the *Load* button to import the model into the application.
- In the **Save the model** panel, the user can save the defined model. Use the *Check model* button to perform basic validation checks before saving. Once the checks are complete, the user can save the entire model by clicking the *Save the model* button. Additionally, the user can link the newly created environment in the *F4F* application to **FLAME GPU 2** by using the *Link the model to FLAME GPU 2* button, giving it a name. In this way, the model will be saved in every directory named **FORGE4FLAME/inst/FLAMEGPU-FORGE4FLAME/resources/f4f** within the user’s home directory. **To run a model in FLAME GPU 2, it must be saved in this location.** The user can create a subdirectory within this path, placing the JSON file and the required RDs downloaded using the *Save the model* button. This process is automated when using the *Link the model to FLAME GPU 2* button.
- In the **References** panel, the user will find the sources and references that were used as a basis for the development of the tool.

4.11 Post Processing

On this page—see Figure 48—users can analyse and visualise simulation results, filter specific simulations of interest, and observe virus evolution in each room, among other aspects. The section is divided into several panels, as described below. Additionally, if *F4F* is run using Docker without Compose, loading the simulation results will not be possible.

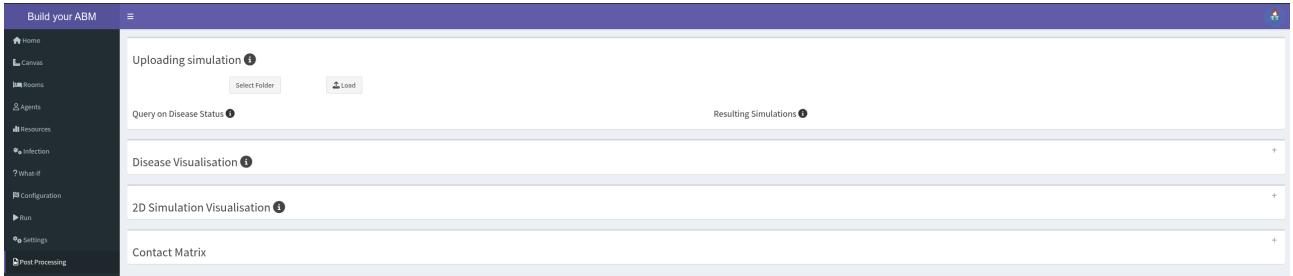


Figure 48: The **Post Processing** page.

4.11.1 Uploading simulation

Here, the user must upload the folder containing all the files and directories generated by the simulation—see Section 3 for more details on where results are saved. The required files, automatically generated by FLAME GPU 2, include `rooms_mapping.txt`, `seed.txt`, and all directories named `seed*`, each containing the following files: `evolution.csv`, `counters.csv`, `AEROSOL.csv`, `CONTACTS_MATRIX.csv`, `CONTACT.csv`, `INFO.csv`, and `AGENT_POSITION_AND_STATUS.csv`. For more details on these files, see Section 6.

Once the directory is chosen, a set of sliders will appear, as shown in Figure 49. These sliders allow users to select a specific subset of simulations—listed on the right side of the figure—by querying them. Users can refine their selection by limiting the simulation days and filtering based on the number of Susceptible, Exposed, Infected, Recovered, and Deceased agents. Based on these filters, the *Resulting Simulations* section will display directories containing simulations that meet the user’s specified ranges. For instance, in Figure 49, only the simulations in which, between days 3 and 6, there were 20-30 susceptible agents, 0 exposed, 3-7 infected, 0-10 recovered, and 0 deaths were selected.

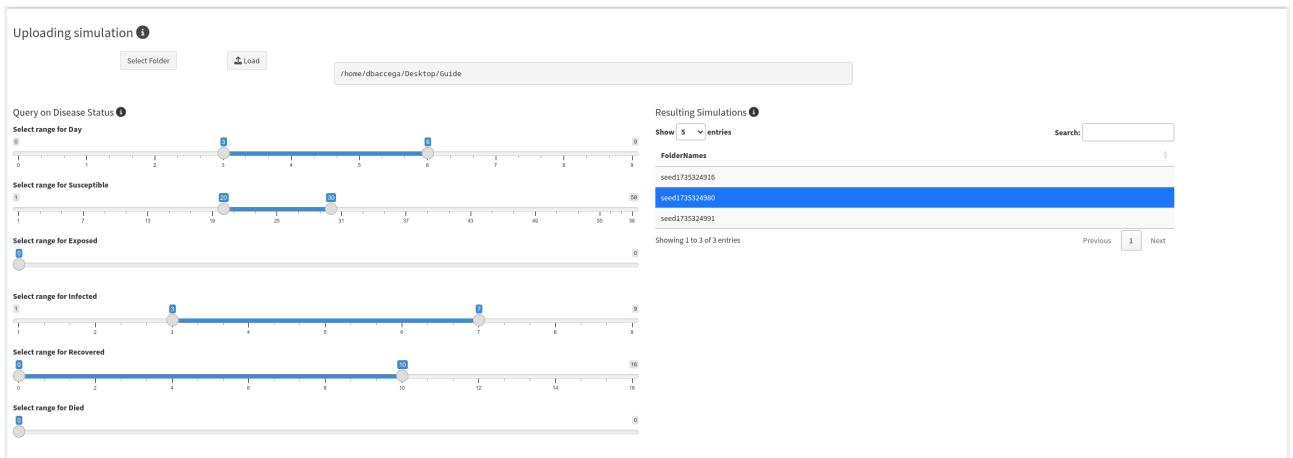


Figure 49: Example of query.

By clicking on each listed directory, the corresponding plots will appear below. As shown in Figure 50, these plots illustrate the daily evolution of the disease, along with additional data such as the number of agents birth or deaths (in the case of agents with a birth rate), the number of agents in quarantine, the number of agents infected outside the environment, and the number of swabs conducted. Additionally, the user can select a specific room and observe how contacts between susceptible and infected agents, as well as the virus concentration, have varied on an hourly basis—a contact between two agents is defined as the situation where they remain close to each other for a certain number of steps without ever separating. Finally, for

each of these plots, the user can display the mean curves along with the areas representing all simulations.

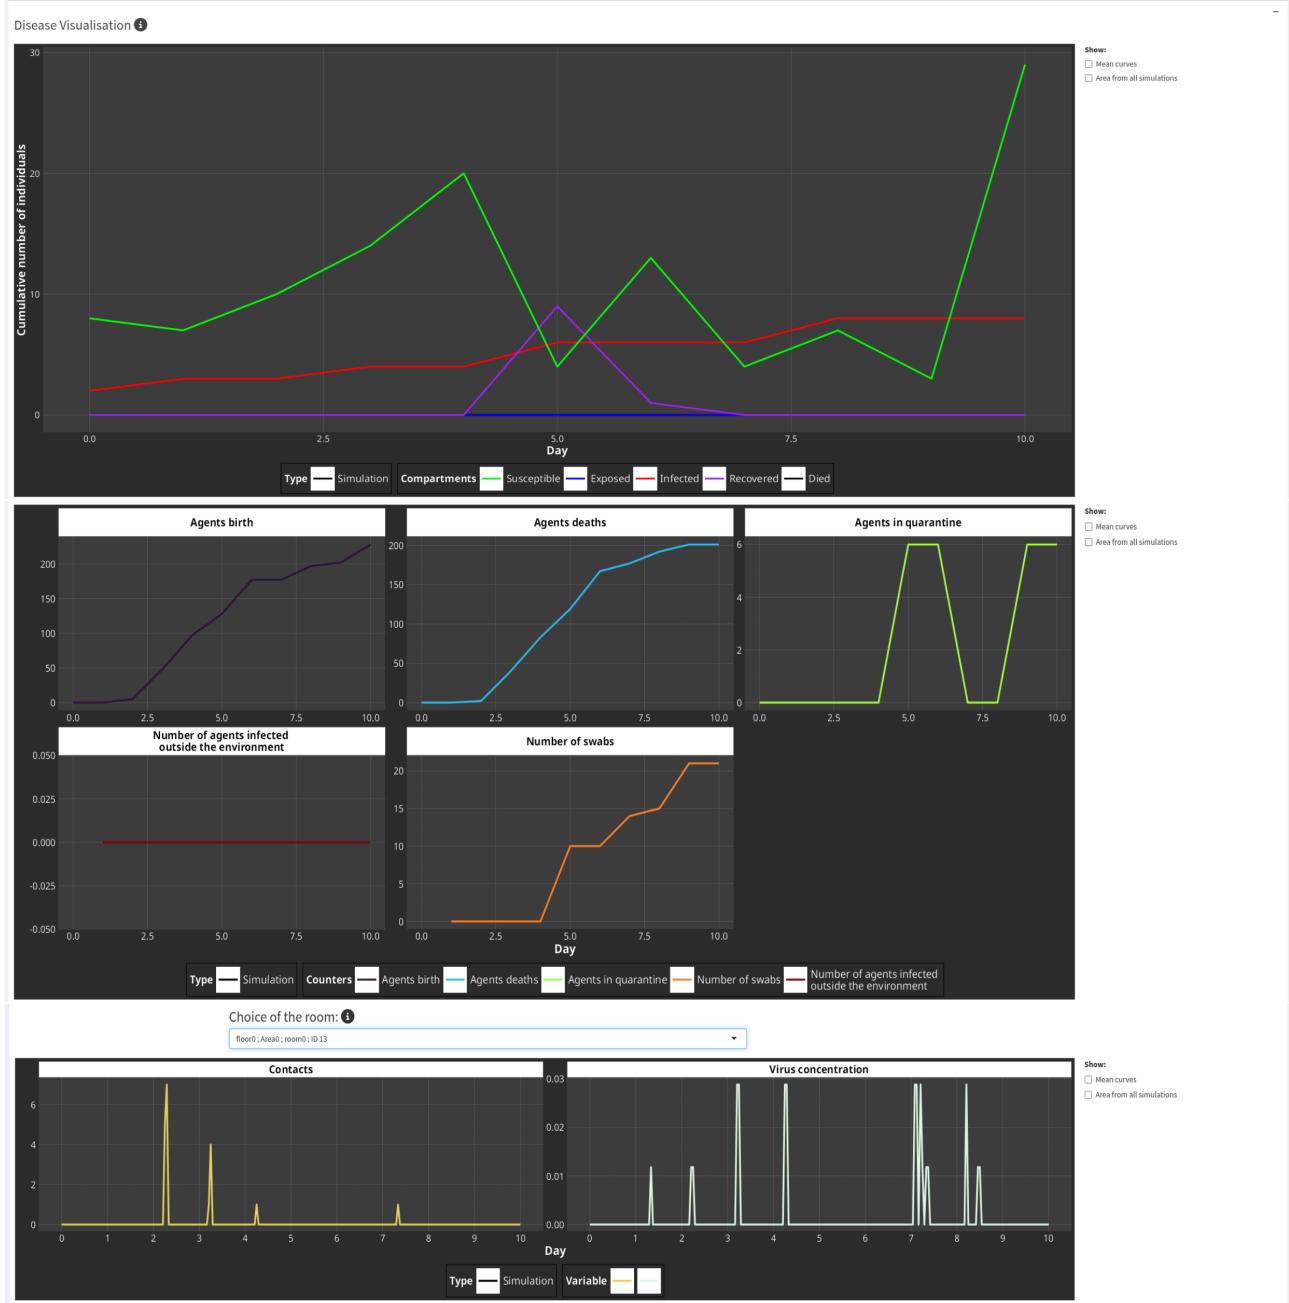


Figure 50: Plots generated after selecting a listed directory.

4.11.2 2D Simulation Visualisation

In this panel—see Figure 51—the user can navigate through two main sections:

- In the *Features* panel, the user can configure various display options:
 - Choose which floors to display (default: all floors).
 - Choose which agent types to display (default: all agent types).
 - Choose room colours based on name, type, area, cumulative number of contacts, current or cumulative virus concentration (default: name).

- Enable additional visual elements, such as room IDs, names, types or areas and agent IDs.

Figure 52 illustrates an example based on the model described in previous sections. Additionally, when an agent type is selected, the user can filter and display specific agents by their IDs—see Figure 53.

- In the *2D Visualisation* panel, the dynamics of the simulation are visualised in a 2D plot. The user can observe agent movements and changes in disease states using a time slider. In particular:

- Different shapes represent different agent types.
- Colours indicate disease states.
- If a shape or colour is missing in the legend, it means no agent with those attributes is currently visible on the selected floor(s).

The animation step can be adjusted via the input field next to the slider. The user can advance the simulation using the *Next* button, which increments the time step by the specified value and updates the 2D visualisation automatically, eliminating the need for manual adjustments. Finally, Figure 54 shows the cumulative virus concentration at the final time.

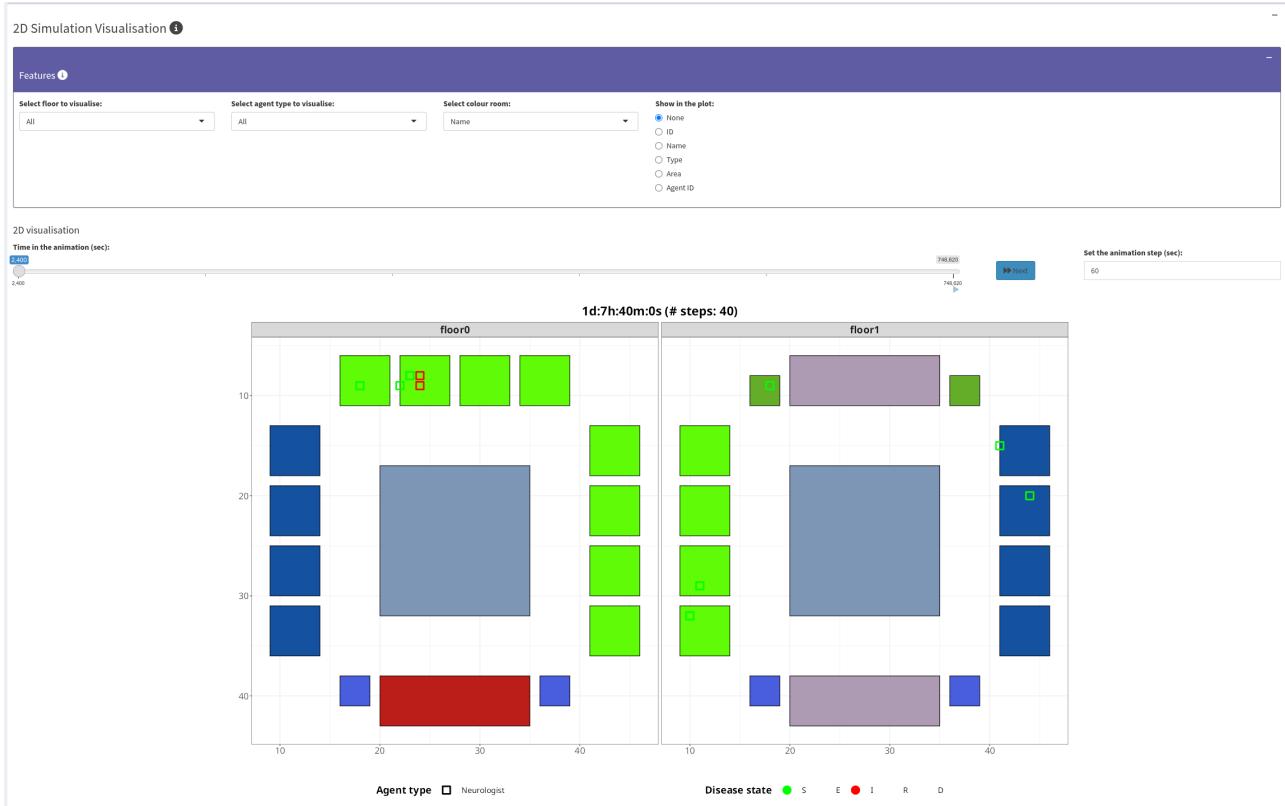


Figure 51: The 2D visualisation section.

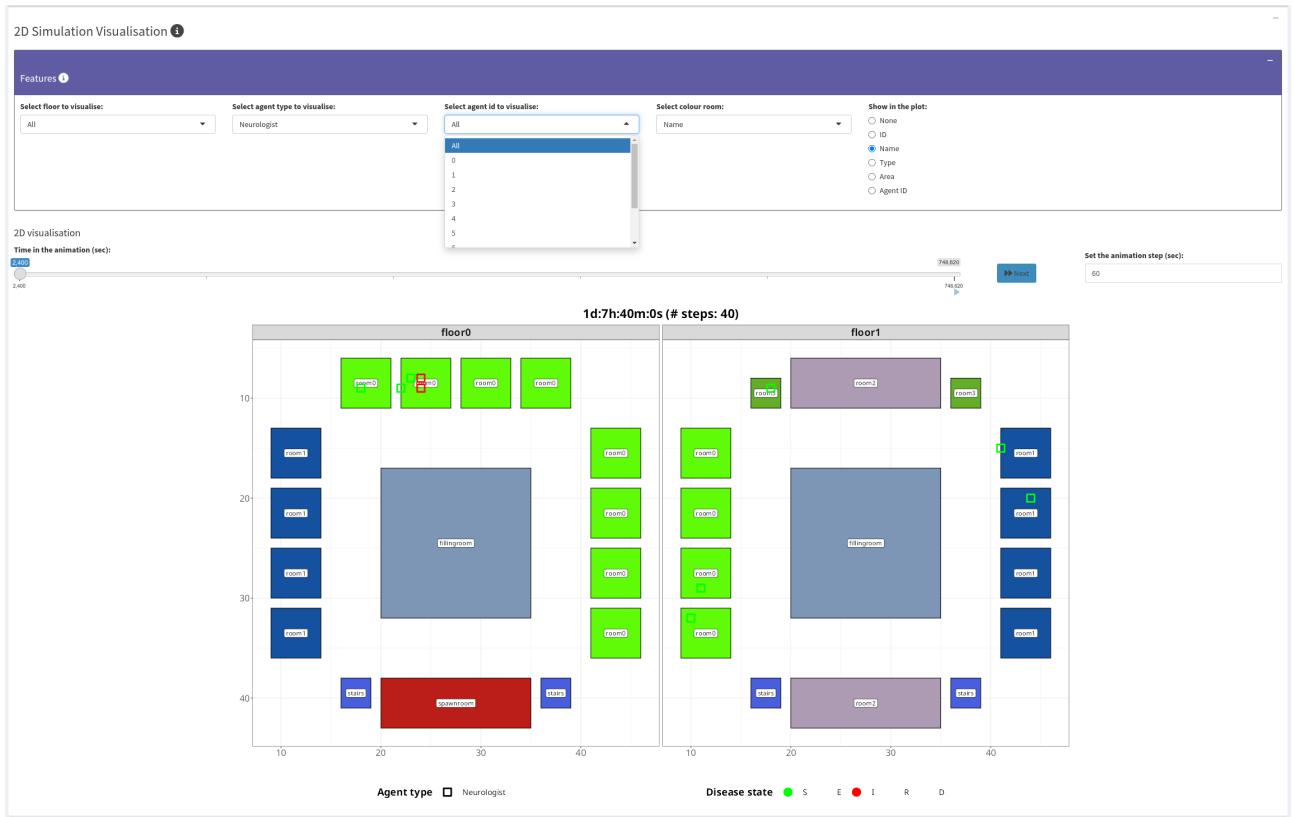


Figure 52: The 2D visualisation section in which rooms' names and only the agent type *Neurologist* are displayed.

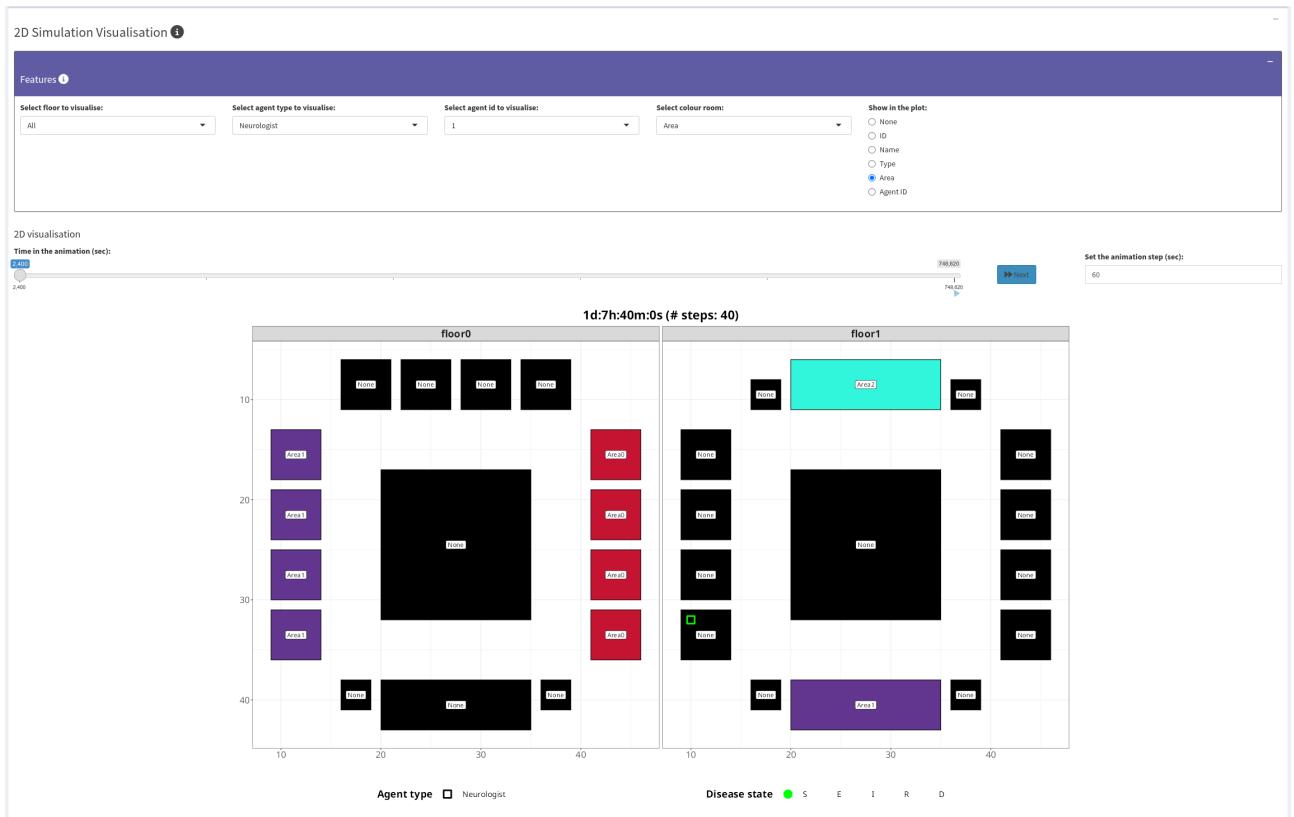


Figure 53: The 2D visualisation section in which rooms' areas and only the agent type *Neurologist* with id 1 are displayed.

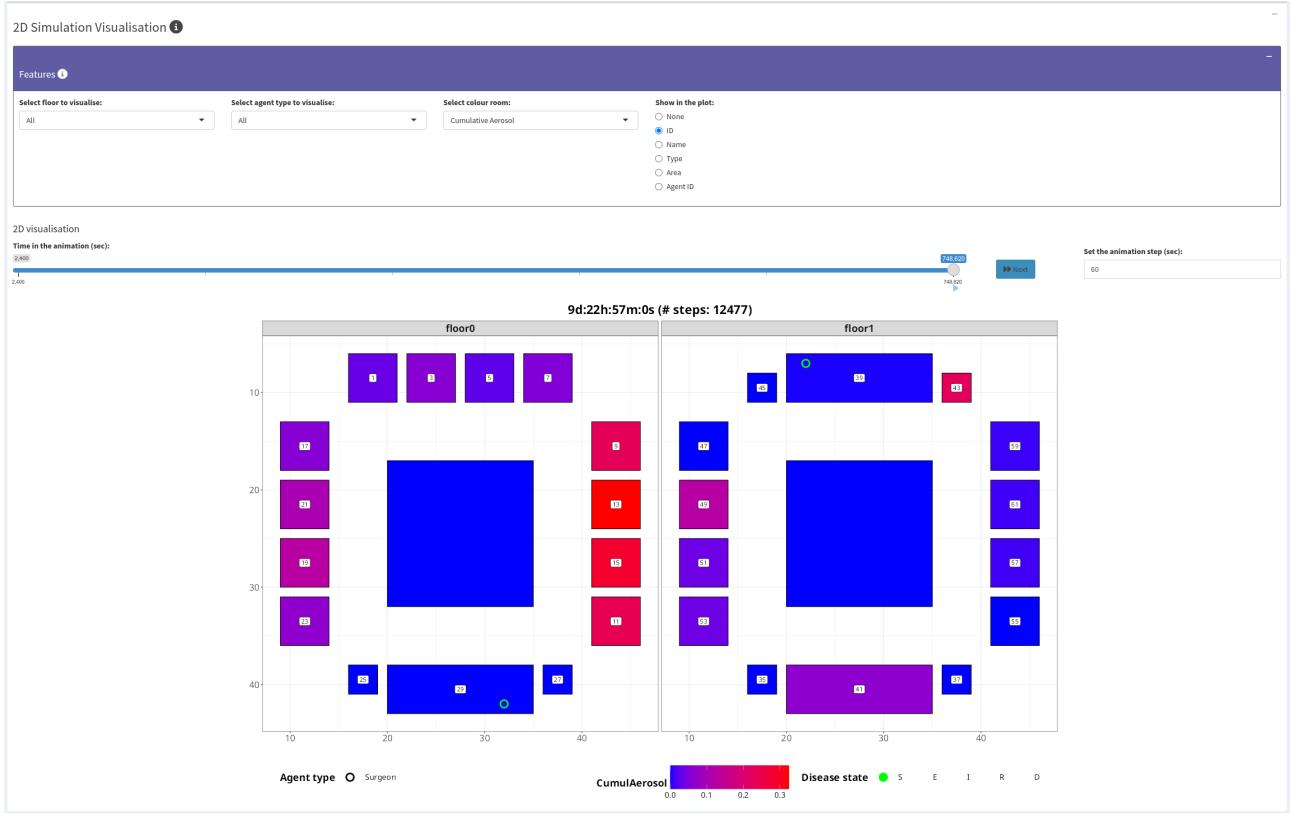


Figure 54: The 2D visualisation section in which the cumulative virus concentration at the final time is shown together with rooms' IDs.

4.11.3 Contact Matrix

In this panel—see Figure 55—the lower triangular contact matrix shows the hourly average number of contacts among the different agent types.

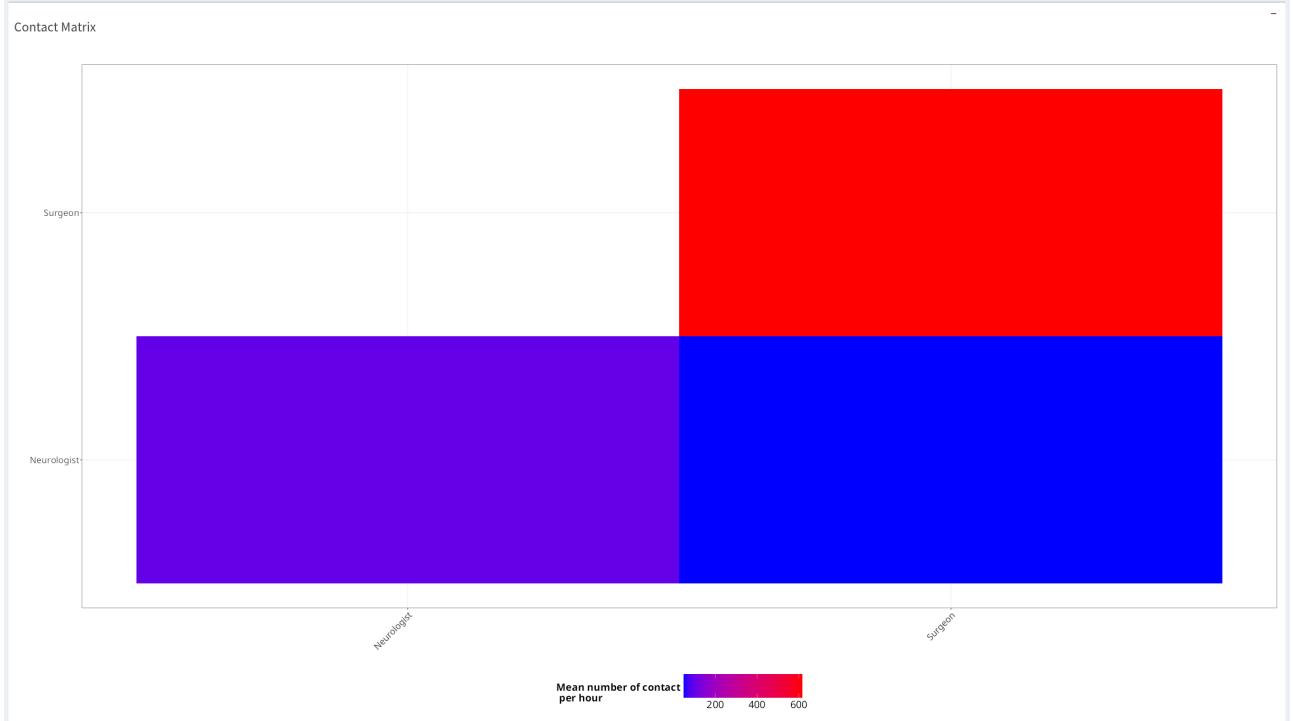


Figure 55: The *Contact Matrix* panel.

5 School

In this section, we provide further details about the generic middle school environment described in the main document. Figures 56, 57, and 58 illustrate the three floors of the school defined using *F4F*. Specifically, *floor0* includes the spawn-room where agents enter the environment, along with four classrooms, two bathrooms—one for students and one for teachers—the gym, the teacher’s room, the principal’s office, the temperature measurement room, the stairs connecting the floors, a couple of filling rooms, and four graph points in the hallway. On *floor1* and *floor2*, there are four classrooms, a student bathroom, stairs, some filling rooms, and two graph points in the hallway.

Figures 59, 60, and 61 illustrate the definition of the agent *student_1A*. Figure 59 presents one of the two determined flows associated with the agent, where the agent spends 300 minutes in classroom 1A, with two 15-minute breaks—one between the second and third classes and another between the fourth and fifth classes. Figure 60 highlights the random events—two random events are possible: going to the bathroom or visiting the principal’s office. Finally, Figure 61 depicts one of the two entry flows linked to the first determined flow, which occurs every day except on Thursday, Saturday, and Sunday.

Figures 62, 63, and 64 illustrate the definition of the agent *teacher_0*. Figure 62 presents one of the three determined flows associated with the agent, where the agent will spend 115 minutes in classroom 1A. Figure 63 highlights the random events—only one possible random event includes going to the bathroom. Finally, Figure 64 depicts one of the three entry flows linked to the first determined flow, which occurs on Monday and Tuesday.

Figure 65 illustrates the infection parameters of the SEIRS compartmental model used in the alert experiments in the main document. Specifically, we set the average incubation and infection periods to 5 days and the average immunisation period to one year and a half.

Figures 66, 67, 68, and 69 present screenshots from the *Post Processing* page considering the *No Countermeasures* scenario introduced in the main paper. The selection criteria for

these simulations required that, between days 20 and 40, the number of agents fell within the following ranges: 100–200 susceptible, 30–60 exposed, 20–40 infected, 0–150 recovered, and 0 deaths. Figure 67 illustrates the trajectories of one of the 56 simulations that meet the query conditions shown in Figure 66. In the 2D visualisation of the school—Figure 68—, the snapshot captures the day 25th of the simulation, depicting students in classrooms, the principal in his office, and a few teachers in the teacher’s room. Additionally, the colour of each room represents the hourly virus concentration. Figure 69 shows the contact matrix among the different types of agents. Finally, we do not include screenshots of the *What-If*, *Configuration*, and *Settings* pages, as users have the flexibility to customize parameters as needed. Furthermore, no restrictions were imposed on room access based on resource availability, as each room was globally assigned 1,000 resources.

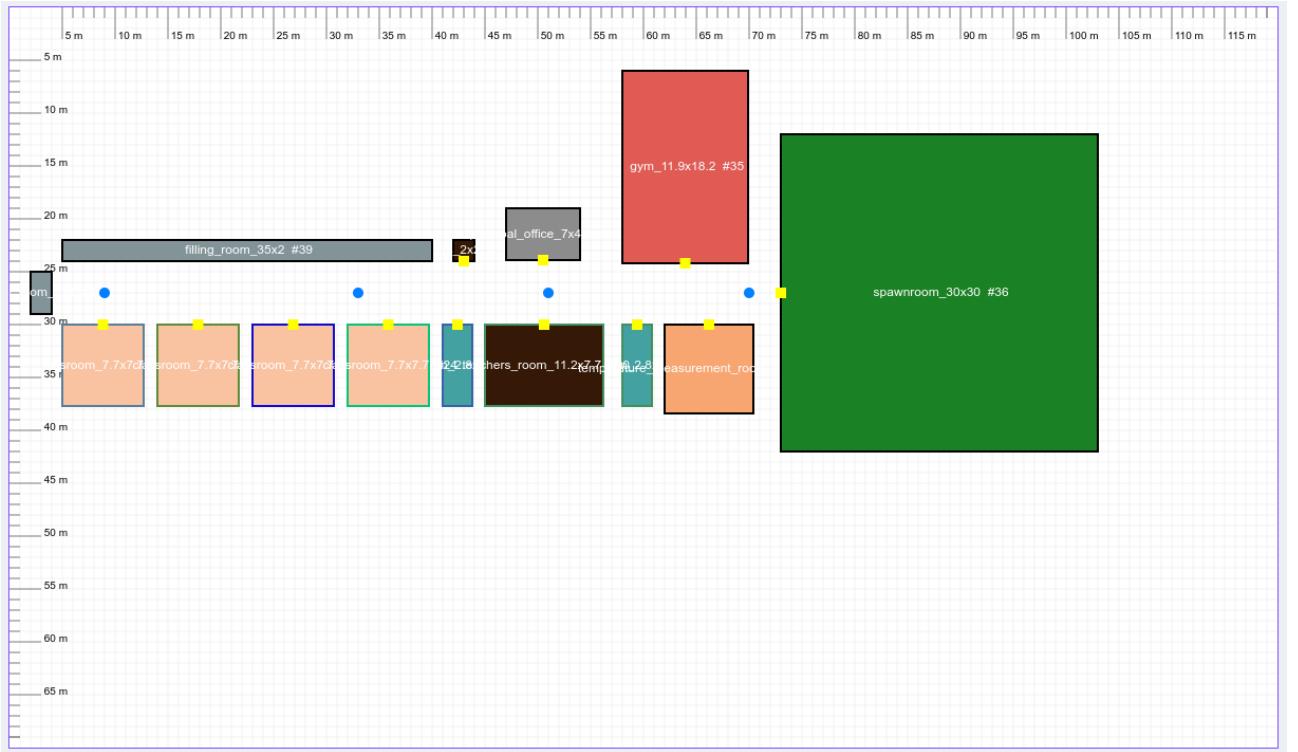


Figure 56: *floor0* of the school environment visualised in the canvas of *F4F*.

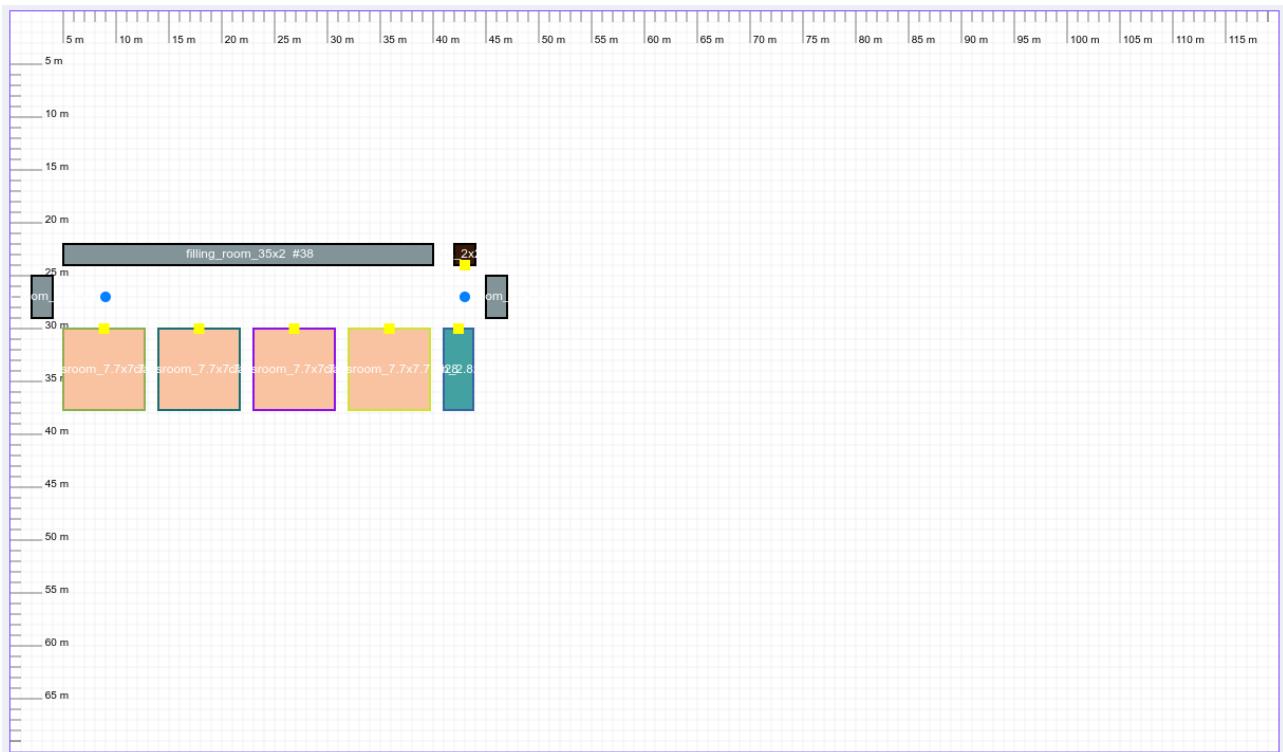


Figure 57: *floor1* of the school environment visualised in the canvas of *F4F*.

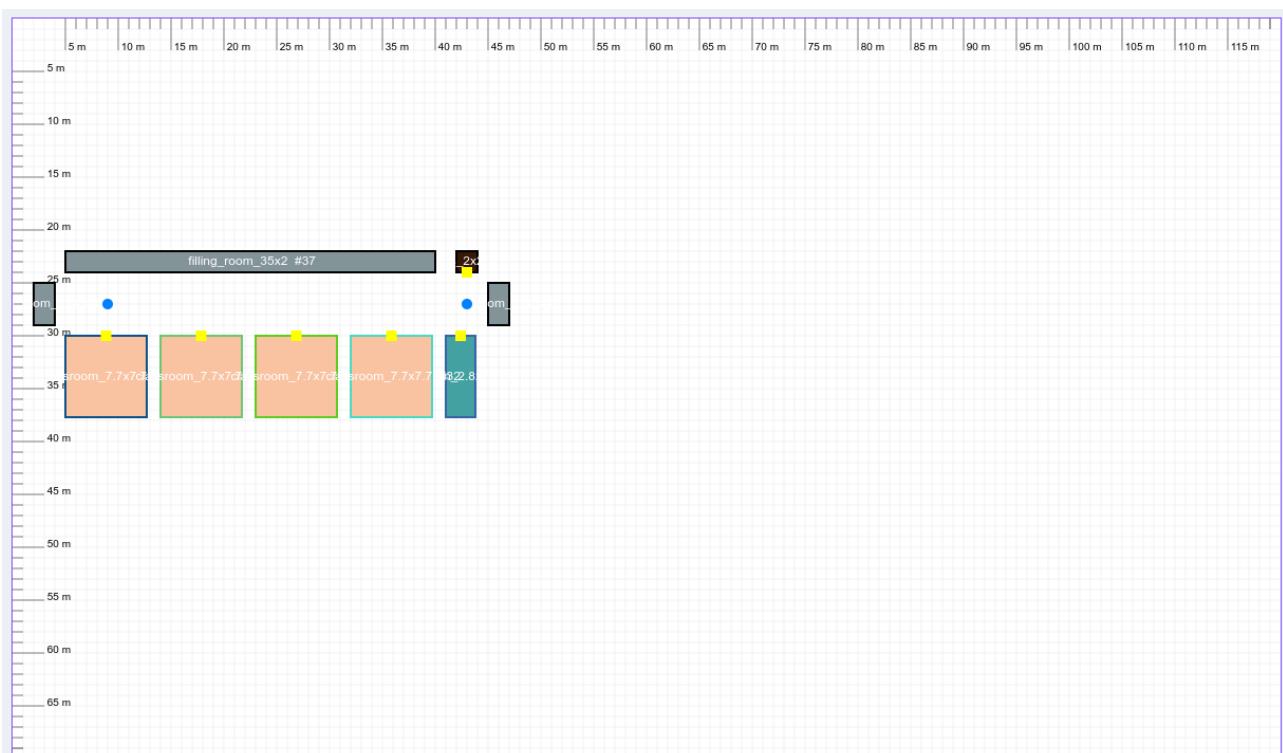


Figure 58: *floor2* of the school environment visualised in the canvas of *F4F*.

Determined flow ①

Type: Activity:

Deterministic Stochastic

Fixed deterministic value: Value

Add room Remove last room

Add flow Remove flow

1 flow 2 flow

Drag the rooms in the desired order

Spawroom-None - Deterministic 35 min - Very Light
 Classroom-1A - Deterministic 50 min - Light
 (2) Classroom-1A - Deterministic 50 min - Light
 Classroom-1A - Deterministic 15 min - Light
 (3) Classroom-1A - Deterministic 50 min - Light
 (4) Classroom-1A - Deterministic 50 min - Light
 (2) Classroom-1A - Deterministic 15 min - Light
 (5) Classroom-1A - Deterministic 50 min - Light
 (6) Classroom-1A - Deterministic 50 min - Light
 Spawroom-None - Deterministic 1 min - Very Light

Figure 59: Determined flow of agent *student_1A*.

Random flow ①

Type: Activity: Weight: Initial time: Ending time:

hh:mm hh:mm

Deterministic Stochastic

Fixed deterministic value: Value

Add room

Click on an event to remove it (except the 'Do nothing' event)

Show 5 entries Search:

Room	Distribution	Activity	Time	Weight	Time Slot
Do nothing	Deterministic	Very Light	0	0.999186	00:00 - 23:59
Bathroom-Student	Exponential	Light	5	8e-04	00:00 - 23:59
Principal_Office-None	Deterministic	Light	50	1.4e-05	00:00 - 23:59
Showing 1 to 3 of 5 entries					

Previous 1 Next

Figure 60: Random flow of agent *student_1A*.

Entry flow

Select type of entrance: Daily Rate Time window

Save time

1 slot 2 slot

Entry time: 07:35

Select Days of the Week

Monday Tuesday Wednesday Thursday Friday Saturday Sunday

Associate with a determined flow: 1 flow

Add slot Remove slot

Figure 61: Entry flow of agent *student_1A*.

Determined flow ①

Type: Activity:

Deterministic Stochastic

Fixed deterministic value: Value

Add room Remove last room

Add flow Remove flow

1 flow 2 flow 3 flow

Drag the rooms in the desired order

Spawroom-None - Deterministic 35 min - Very Light
 Classroom-1A - Deterministic 115 min - Hard
 Spawroom-None - Deterministic 1 min - Very Light

Figure 62: Determined flow of agent *teacher_0*.

Random flow ①

Type:	Activity:	Weight:	Initial time:	hh:mm	Ending time:	hh:mm
Deterministic	Stochastic					
Fixed deterministic value:		<input type="text"/>				
Value:						
Click on an event to remove it (except the 'Do nothing' event)						
Show 5 entries						
Room	Distribution	Activity	Time	Weight	Time Slot	
Do nothing	Deterministic	Very Light	0	0.999200000000001	00:00 - 23:59	
Bathroom-Teacher	Exponential	Quite Hard	5	0.0008	00:00 - 23:59	
Showing 1 to 2 of 2 entries						
						Previous 1 Next

Figure 63: Random flow of agent *teacher_0*.

Entry flow

Select type of entrance:	<input type="radio"/> Daily Rate <input checked="" type="radio"/> Time window	
Save time	<input type="button"/>	
1 slot	2 slot	3 slot
Entry time:	<input type="text"/> 07:35	
Associate with a determined flow:	<input type="button"/> 1 flow	
Select Days of the Week	<input checked="" type="checkbox"/> Monday <input checked="" type="checkbox"/> Tuesday <input type="checkbox"/> Wednesday <input type="checkbox"/> Thursday <input type="checkbox"/> Friday <input type="checkbox"/> Saturday <input type="checkbox"/> Sunday	
Add slot	<input type="button"/> Remove slot	

Figure 64: Entry flow of agent *teacher_0*.

Model: SEIRS

Disease model: SEIRS. Beta(contact): 0.024, Beta(aerosol): 410, Gamma: 0.2 (Exponential), Alpha: 0.2 (Exponential), Nu: 0.001826484 (Exponential)

Description:
The SEIRS model is an extension of the SEIR model that incorporates the concept of temporary immunity, resulting in a cycle where individuals can return to susceptibility. It includes four compartments. The first compartment, susceptible (S), represents individuals who are vulnerable to the disease. When these individuals come into contact with infected individuals or share a confined space with them, they can contract the disease and transition to the exposed compartment. The second compartment, exposed (E), includes individuals who have been infected but are in an incubation period during which they cannot yet transmit the disease. The third compartment, infected (I), comprises individuals who are actively infected and capable of spreading the disease to susceptible individuals. The fourth compartment, recovered (R), consists of individuals who have overcome the infection and gained immunity. In contrast to the SEIR model, the SEIRS model assumes that immunity is temporary. After a certain period, individuals in the recovered compartment lose their immunity and return to the susceptible compartment, making the SEIRS model particularly suitable for studying diseases where immunity wanes over time.

β _{contact} =	0.024
β _{aerosol} =	410
γ =	<input type="button"/>
Deterministic	<input type="radio"/> Stochastic
Distribution:	<input type="button"/> Exponential
Value:	0.2
α =	<input type="button"/>
Deterministic	<input type="radio"/> Stochastic
Distribution:	<input type="button"/> Exponential
Value:	0.2
ν =	<input type="button"/>
Deterministic	<input type="radio"/> Stochastic
Distribution:	<input type="button"/> Exponential
Value:	0.001826484

Figure 65: Infection parameters used in the alert experiments.

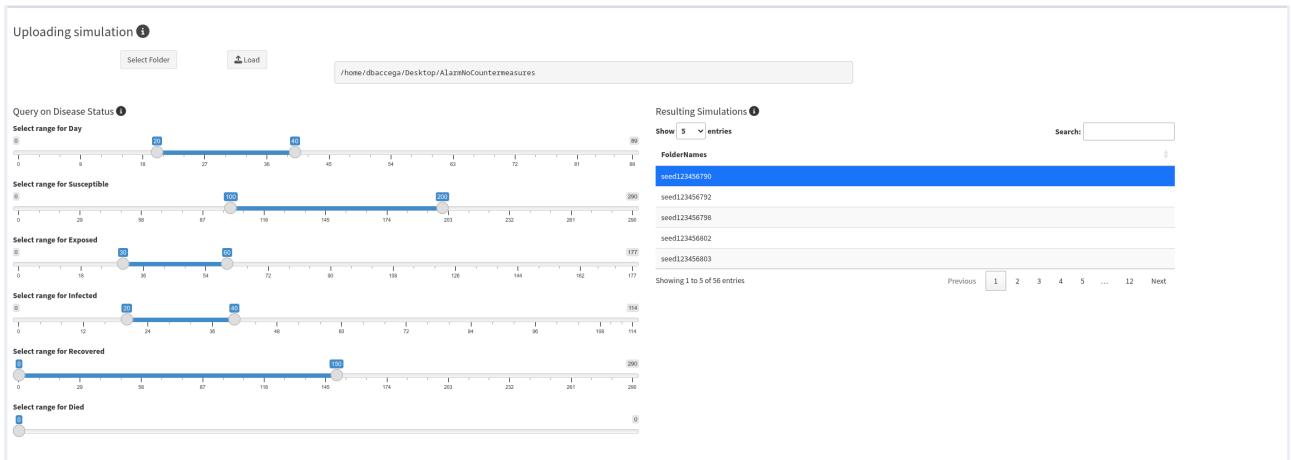


Figure 66: Example of query on the results obtained using the school model on the *No Countermeasures* scenario.

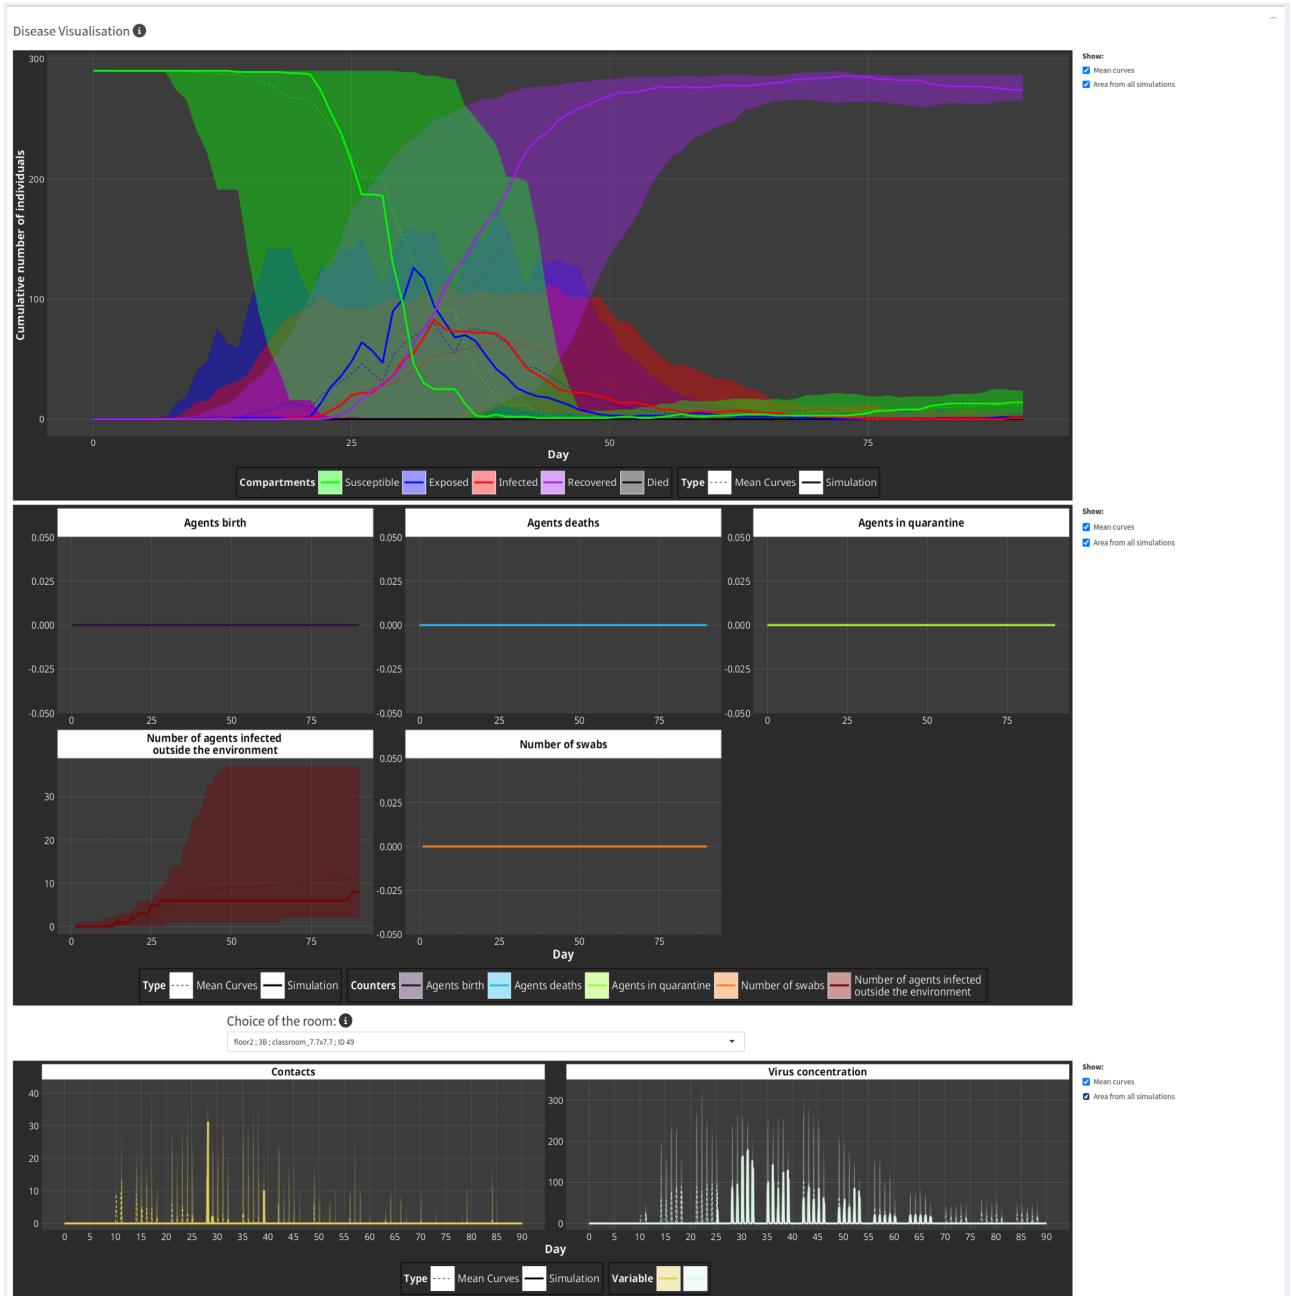


Figure 67: Plots generated after selecting a listed directory using the school model on the *No Countermeasures* scenario (with means and areas).



Figure 68: 2D visualisation of the school where the colour of each room represents the hourly virus concentration on the *No Countermeasures* scenario.

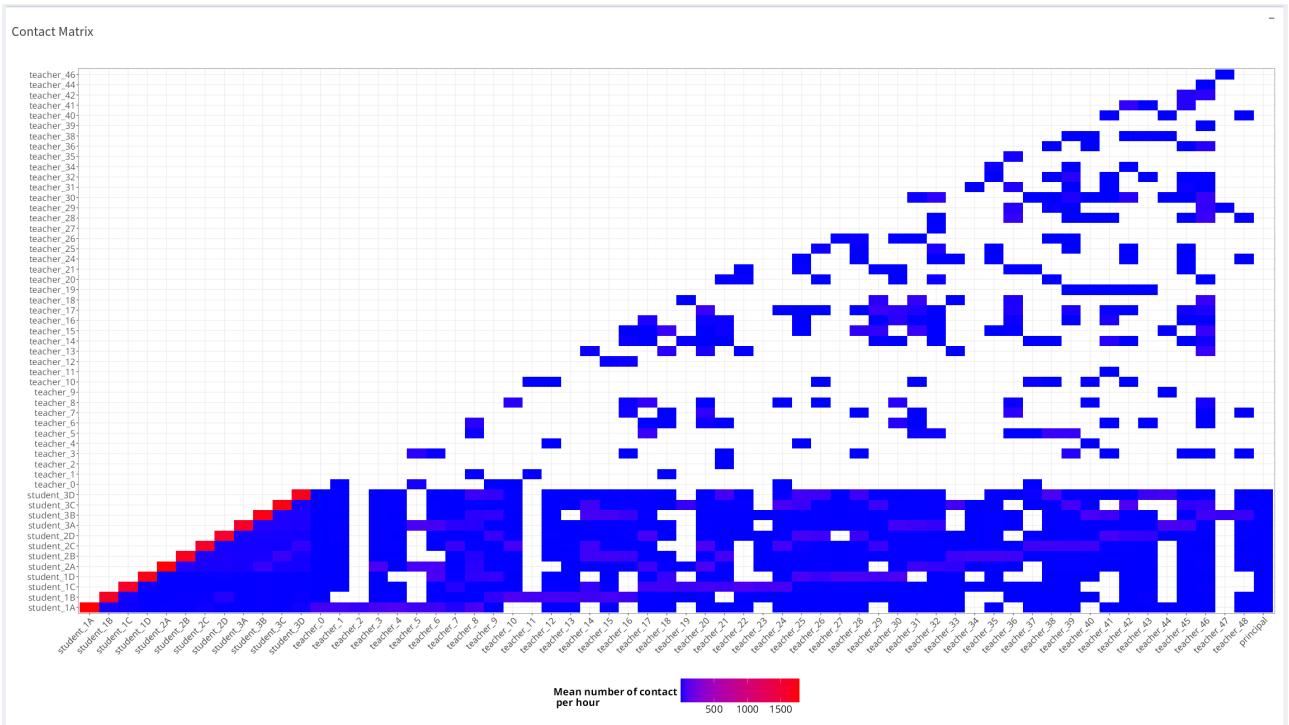


Figure 69: Contact matrix of the school model on the *No Countermeasures* scenario.

6 Link between F4F and FLAME GPU 2

As said in the main document, the output of *F4F*—a JSON file including the whole model—is linked to FLAME GPU 2 through a Python script that parses the JSON file to generate the

necessary FLAME GPU 2 code and XML files based on the model defined in *F4F*. Specifically, it produces header files (.h) containing definitions for dynamic simulation parameters such as the chosen disease model, the environment's dimensions, the initial time, and the number of days to simulate. It also generates several XML files, which are later read by FLAME GPU 2. These XML files include data for rooms, agents, environmental, and macro-environmental variables

In the main document, we described the movement graph. This graph is automatically constructed, with a vertex created for each room centre, door, and user-defined graph point (particularly useful for placing vertices near corridor corners, improving pathfinding accuracy). Two vertices are connected by an edge if and only if there are no walls obstructing the line of sight between them, or (if they represent stairs) if they are approximately aligned in position across different floors. The graph is stored in FLAME GPU 2's environmental and macro-environmental variables, allowing agents to navigate the space efficiently. When movement is required, agents rely on the A* algorithm, combined with a cost-surface model, to compute the most efficient path between a source and a destination vertex. For instance, Figure 70 illustrates all possible edges of the graph for *floor0* in the model defined in Section 4.

When running the FLAME GPU 2 simulation with the visualisation enabled—without Docker, see Section 3 for more details—the environment created using *F4F*—as defined in Section 4.2—will be displayed, as shown in Figure 71. Moreover, in Figure 72 a 3D visualisation of the school environment is shown.

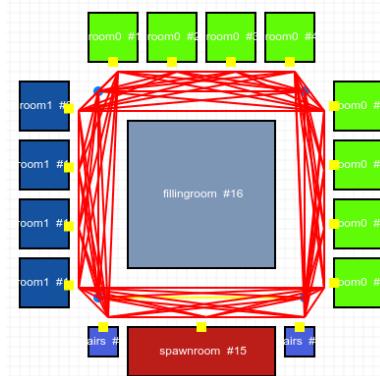


Figure 70: All possible edges of the graph for *floor0* in the model defined in Section 4

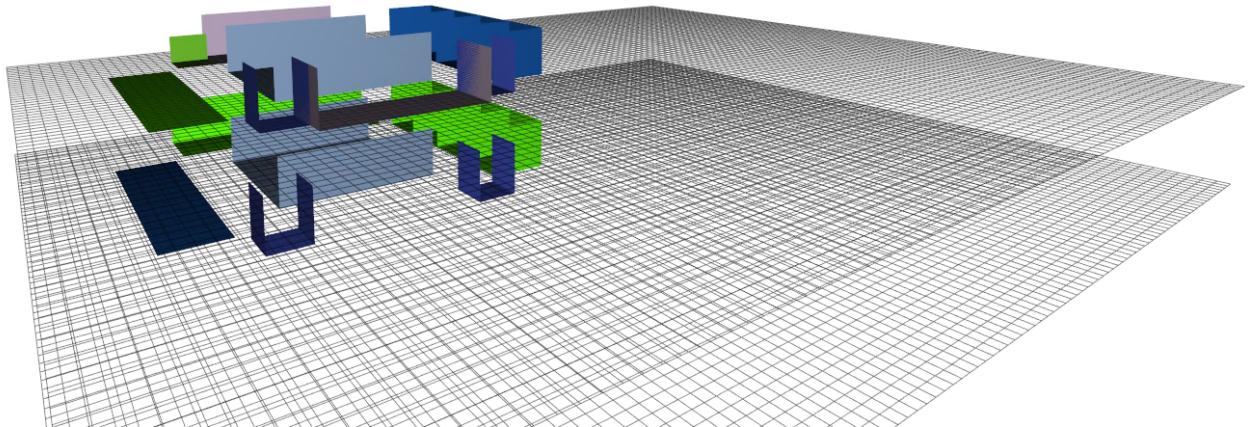


Figure 71: 3D visualisation in FLAME GPU 2 of the example depicted in Section 4.

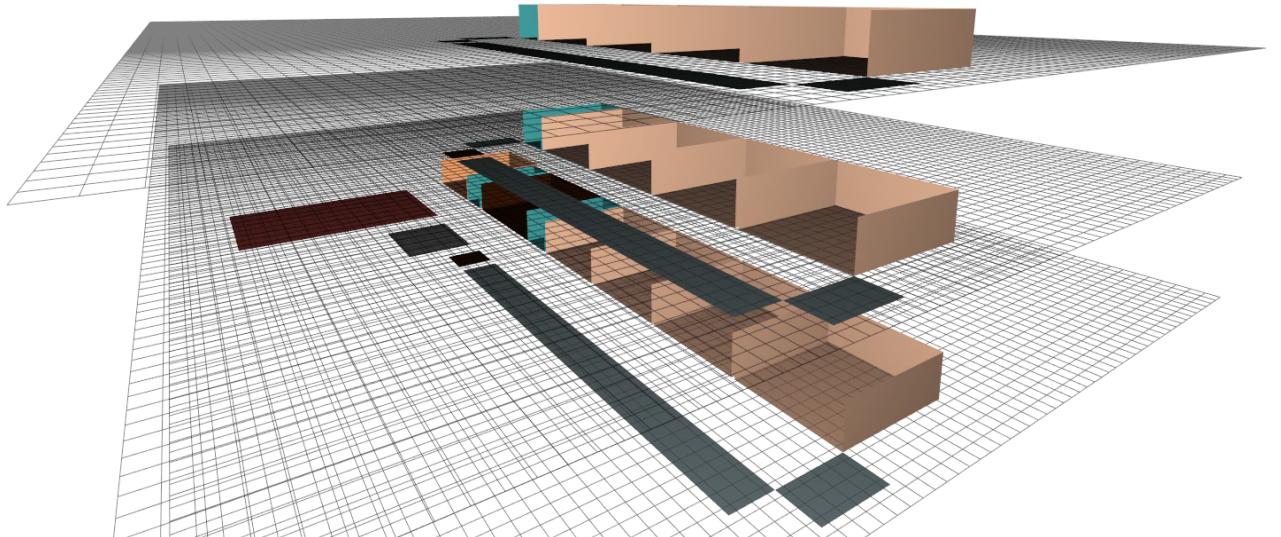


Figure 72: 3D visualisation in FLAME GPU 2 of the school depicted in Section 5 and used in the main document.

In the main document and Section 3, we outline the ways users can run *F4F* and FLAME GPU 2, explaining how users can effectively run simulations. When FLAME GPU 2 completes the simulation runs, it generates CSV files containing various information required to post-analyse the simulated scenario (e.g., disease evolution, contact matrix, virus diffusion within the environment, and agents' movements). This information is stored in separate files—one for disease evolution, one for the contact matrix, and so on. These files can be passed to *F4F* for post-analysis and visualisation of disease dynamics in the modelled system—see Section 4.11.

Specifically, each folder containing the model's results includes:

- `rooms_mapping.txt` stores the mapping between each room's ID in the graph and the corresponding position of its centre;
- `seed.txt` stores the simulation seed specified using *F4F*;

- `WHOLEmodel.json` and `WHOLEmodel.RDs` store the *F4F* model;
- various directories named `seed*`—there are as many directories as there are simulation runs—, each containing the following files:
 - `evolution.csv` stores the daily evolution of the disease; the file contains six columns: *Day, Susceptible, Exposed, Infected, Recovered, and Died*;
 - `counters.csv` stores the daily evolution of some interesting pre-defined counters—already introduced in Section 4.11; the file contains five columns: *Day, COUNTERS_CREATED_AGENTS_WITH_RATE, COUNTERS_KILLED_AGENTS_WITH_RATE, AGENTS_IN_QUARANTINE, SWABS, and NUM_INFECTED_OUTSIDE*;
 - `AEROSOL.csv` stores the hourly evolution of virus concentration in each room—entries are only logged when the virus concentration is greater than zero; the file contains three columns: the first represents the FLAME GPU 2 step—i.e., simulation time—, the second indicates the virus concentration, and the third corresponds to the room’s ID in the graph;
 - `CONTACTS_MATRIX.csv` stores the hourly lower triangular contact matrix—entries are logged only if the number of contacts in the last hour is greater than zero; the file contains four columns: the first represents the FLAME GPU 2 step, the second and third indicate the agent types, and the fourth records the number of contacts in the last hour between those agent types;
 - `CONTACT.csv` stores interactions between susceptible and infected agents; it contains four columns: the first indicates the FLAME GPU 2 step, the second and third specify the types of the susceptible and infected agents, respectively, and the fourth provides the ID of the room where the contact occurred; if agents remain in the contact area from step t_1 to t_2 —where $t_1 < t_2$ —, the user will find $t_2 - t_1$ rows in the file, corresponding to the time steps $t_1, t_1 + 1, \dots, t_2$;
 - `AGENT_POSITION_AND_STATUS.csv` stores the position of each agent whenever it changes from the previous FLAME GPU 2 step; the file contains seven columns: the first indicates the FLAME GPU 2 step, the second represents the agent’s ID, the third specifies the agent type, the fourth, fifth, and sixth provide the agent’s position on the x, y, and z axes, respectively, and the seventh denotes the agent’s disease state.
 - `INFO.csv` contains information about the simulation. Its columns are formatted as follows: the first column indicates the FLAME GPU 2 step, and the second column is a string.
 - `host_rng_state.txt` stores the state of the host’s random number generator at the end of the simulation.

References

- [1] Marco Aldinucci et al. “HPC4AI: an AI-on-demand federated platform endeavour”. In: *Proceedings of the 15th ACM International Conference on Computing Frontiers*. 2018, pp. 279–286. DOI: <https://doi.org/10.1145/3203217.3205340>.
- [2] D Attali. “Colourpicker: A colour picker tool for shiny and for selecting colours in plots”. In: *R package version 1* (2017). URL: <https://cran.r-project.org/web/packages/colourpicker/>.

- [3] D Attali, T Edwards, and Z Wang. “shinyalert: Easily create pretty popup messages (modals) in ‘shiny’”. In: *R package version 2.0* (2018). URL: <https://cran.r-project.org/web/packages/shinyalert/>.
- [4] Dean Attali. “Shinyjs: Easily improve the user experience of your shiny apps in seconds”. In: *R package version 2.0* (2020). URL: <https://cran.r-project.org/web/packages/shinyjs/>.
- [5] Eric Bailey. “shinyBS: Twitter bootstrap components for shiny”. In: *R package version 0.61* (2015). URL: <https://cran.r-project.org/web/packages/shinyBS/>.
- [6] Winston Chang et al. “Package ‘shiny’”. In: *See http://citeseervx.ist.psu.edu/viewdoc/-download* (2015). URL: <https://cran.r-project.org/web/packages/shiny/>.
- [7] Winston Chang et al. “Package ‘shinydashboard’”. In: (2022). URL: <https://cran.r-project.org/web/packages/shinydashboard/>.
- [8] Winston Chang et al. *shiny: Web Application Framework for R*. R package version 1.10.0, <https://github.com/rstudio/shiny>. 2024. URL: <https://shiny.posit.co/>.
- [9] Winston Chang et al. “shinythemes: Themes for Shiny”. In: *R package version 1.2* (2018). URL: <https://cran.r-project.org/web/packages/shinythemes/>.
- [10] Gábor Csárdi, Kuba Podgóński, and Rich Geldreich. *zip: Cross-Platform “zip” Compression*. 2020. URL: <https://cran.r-project.org/web/packages/zip/>.
- [11] Gábor Csárdi et al. *remotes: R Package Installation from Remote Repositories, Including ‘GitHub’*. R package version 2.5.0, <https://github.com/r-lib/remotes#readme>. 2024. URL: <https://remotes.r-lib.org>.
- [12] Savvas Gkantonas et al. “airborne.cam: A risk calculator of SARS-CoV-2 aerosol transmission under well-mixed ventilation conditions”. In: (2021). DOI: <https://doi.org/10.17863/CAM.72192>.
- [13] Jim Hester and Jennifer Bryan. “Glue: Interpreted string literals”. In: *R package version 1.2* (2020). URL: <https://cran.r-project.org/web/packages/glue/>.
- [14] Nicolas Hoertel et al. “A stochastic agent-based model of the SARS-CoV-2 epidemic in France”. In: *Nature medicine* 26.9 (2020), pp. 1417–1421. DOI: <https://doi.org/10.1038/s41591-020-1001-6>.
- [15] Fanny Meyer and Victor Perrier. *shinybusy: Busy Indicator for “Shiny”Applications*. 2020. URL: <https://cran.r-project.org/web/packages/shinybusy/>.
- [16] Jeroen Ooms, Duncan T Lang, and Lloyd Hilaiel. “jsonlite: A robust, high performance JSON parser and generator for R”. In: *R package version 1* (2018). DOI: <https://cran.r-project.org/web/packages/jsonlite/>.
- [17] David Ornelles. *EBImage Extra: Helper functions for EBImage*. Version 0.0.0.2. License: GPL-3. 2019. DOI: <https://orcid.org/0000-0002-5151-8382>. URL: <https://github.com/ornelles/EBImageExtra>.
- [18] Grégoire Pau et al. “EBImage—an R package for image processing with applications to cellular phenotypes”. In: *Bioinformatics* 26.7 (2010), pp. 979–981. DOI: <https://doi.org/10.1093/bioinformatics/btq046>.
- [19] Thomas Lin Pedersen et al. “shinyFiles: a server-side file system viewer for Shiny”. In: *R package version 0.6.2* (2016).

- [20] Zhe Peng et al. “Practical indicators for risk of airborne transmission in shared indoor environments and their application to COVID-19 outbreaks”. In: *Environmental science & technology* 56.2 (2022), pp. 1125–1137. DOI: <https://doi.org/10.1021/acs.est.1c06531>.
- [21] Victor Perrier, Fanny Meyer, David Granjon, et al. “shinyWidgets: Custom inputs widgets for Shiny”. In: *R package version 0.4_7* (2019). URL: <https://cran.r-project.org/web/packages/shinyWidgets/>.
- [22] RStudio and Inc. “htmltools: Tools for HTML”. In: *R package version 3* (2017), p. 6.
- [23] Kristin Tolksdorf et al. “Influenza-associated pneumonia as reference to assess seriousness of coronavirus disease (COVID-19)”. In: *Eurosurveillance* 25.11 (2020), p. 2000258. DOI: <https://doi.org/10.2807/1560-7917.es.2020.25.11.2000258>.
- [24] A de Vries, B Schloerke, and K Russell. *sortable: Drag-and-Drop in ‘shiny’Apps with ‘SortableJS’*. *R package version 0.5.0*. 2024. URL: <https://cran.r-project.org/web/packages/sortable/>.
- [25] Hadley Wickham. *ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York, 2016. ISBN: 978-3-319-24277-4. URL: <https://ggplot2.tidyverse.org>.
- [26] Hadley Wickham. “Package ‘tidyR’”. In: *Easily Tidy Data with ‘spread’and ‘gather ()’Functions* (2017). URL: <https://cran.r-project.org/web/packages/tidyr/>.
- [27] Hadley Wickham. “stringr: modern, consistent string processing”. In: *The R Journal* 2.2 (2010), pp. 38–40. DOI: <https://doi.org/10.32614/RJ-2010-012>. URL: <https://doi.org/10.32614/RJ-2010-012>.
- [28] Hadley Wickham et al. *devtools: Tools to Make Developing R Packages Easier*. <https://devtools.r-lib.org/>, <https://github.com/r-lib/devtools>. 2022.
- [29] Hadley Wickham et al. *dplyr: A Grammar of Data Manipulation*. R package version 1.1.4. 2023. URL: <https://CRAN.R-project.org/package=dplyr>.
- [30] Hadley Wickham et al. “Package ‘readr’”. In: *Read Rectangular Text Data. Available online: https://cran.r-project.org/web/packages/readr/readr.pdf (accessed on 23 August 2023)* (2024). URL: <https://cran.r-project.org/web/packages/readr/>.
- [31] Yihui Xie, Joe Cheng, Xianying Tan, et al. “DT: a wrapper of the JavaScript library ‘DataTables’”. In: *R package version 0.4* (2018). URL: <https://cran.r-project.org/web/packages/DT/>.